

Optimization 2

Master Data Science

S. Gaïffas



- 1 Stochastic optimization
 - Stochastic Gradient Descent (SGD)
 - Beyond SGD

- 2 Supervised learning recipes (bis)
 - File indexing
 - Lazy-updates
 - Cross-validation
 - Warm starting

- 1 Stochastic optimization
 - Stochastic Gradient Descent (SGD)
 - Beyond SGD

- 2 Supervised learning recipes (bis)
 - File indexing
 - Lazy-updates
 - Cross-validation
 - Warm starting

We want to minimize

$$f(\theta) = \frac{1}{n} \sum_{i=1}^n f_i(\theta)$$

where $f_i(\theta) = \ell(y_i, \langle x_i, \theta \rangle) + \frac{\lambda}{2} \|\theta\|_2^2$

- Ridge Regression
- Ridge Logistic Regression
- etc.

Full (Batch) Gradient Descent

$$\theta^k \leftarrow \theta^{k-1} - \eta_k \nabla f(\theta^{k-1})$$

You've seen:

- If f convex and L -smooth then for $\eta_k = 1/L$

$$f(\theta^k) - f(\theta_*) \leq \frac{2L\|\theta^0 - \theta_*\|}{k+1} \quad (1)$$

where $\theta_* \in \operatorname{argmin}_{\theta} f(\theta)$

- Acceleration (Nesterov, Fista): rate improvement $O(1/k^2)$
- Linesearch

If f is also μ -strongly convex, then linear convergence

$$f(\theta^k) - f(\theta_*) \leq \left(1 - \frac{L}{\mu}\right)^k (f(\theta_0) - f(\theta_*)) \quad (2)$$

What if n (and d) is large?

- Each iteration of a full gradient method has complexity $O(nd)$
- I can't put $n \times d$ floats (32 or 64 bits) in my memory

Size of big data makes a modern computer look old: go back to “old” algorithms

- Idea: in machine learning, objective functions are averages of losses

If I choose uniformly at random $I \in \{1, \dots, n\}$, then

$$\mathbb{E}[\nabla f_I(\theta)] = \frac{1}{n} \sum_{i=1}^n \nabla f_i(\theta) = \nabla f(\theta)$$

- $\nabla f_I(\theta)$ is an *unbiased* but very noisy estimate of the full gradient $\nabla f(\theta)$
- Computation of $\nabla f_I(\theta)$ only requires the I -th line of data ($O(d)$ and smaller for sparse data, see next)

Stochastic Gradient Descent (SGD) algorithm (Robbins and Monro 1951)

- At each iteration, load a line of data chosen randomly (requires an index for fast random access on the hard drive)
- Compute gradient for this line of data
- Do a descent step, using this gradient, and repeat

Stochastic Gradient Descent (SGD)

- **Input:** starting point θ^0 , sequence of learning rates $\{\eta_t\}_{t \geq 0}$
- For $t = 1, 2, \dots$ until *convergence* do
 - Pick at random (uniformly) i in $\{1, \dots, n\}$
 - Put

$$\theta^t = \theta^{t-1} - \eta_t \nabla f_i(\theta^{t-1})$$

- **Return** last θ^t
- Each iteration has complexity $O(d)$ instead of $O(nd)$ for full gradient methods
- Possible to reduce this to $O(s)$ when features are s -sparse using **lazy-updates** (more on this later)

- Note that if i is chosen uniformly at random in $\{1, \dots, n\}$

$$\mathbb{E}[\nabla f_i(\theta^{t-1}) | \mathcal{F}_{t-1}] = \frac{1}{n} \sum_{i'=1}^n \nabla f_{i'}(\theta^{t-1}) = \nabla f(\theta^{t-1})$$

where \mathcal{F}_t = information until iteration t (relative to random sampling of indexes)

- Namely, SGD uses very noisy unbiased estimations of the full gradient
- Learning rate η_t usually chosen as $\eta_t \approx Ct^{-\alpha}$ with $\alpha \in [1/2, 1]$.
- Linesearch: η_t is a valid learning rate if

$$f_i(\theta^t - \eta_t \nabla f_i(\theta^t)) \leq f_i(\theta^t) - \frac{\eta_t}{2} \|\nabla f_i(\theta^t)\|_2^2$$

Polyak-Ruppert averaging (ASGD)

- Use SGD iterates $\{\theta^t\}$ but return $\bar{\theta}^t = \frac{1}{t} \sum_{t'=1}^t \theta^{t'}$
- Computed “online”

$$\bar{\theta}^t \leftarrow \frac{1}{t} \theta^{t-1} + \frac{t-1}{t} \bar{\theta}^{t-1} \quad (3)$$

- Leads to better results, cf:

<http://leon.bottou.org/projects/sgd>

Theoretical knowledge on SGD. Typical assumptions

- Each f_i is L -smooth (gradient L -lipshitz)
- f is μ -strongly convex
- Non strongly convex: rate $O(1/\sqrt{t})$ for ASGD with $\eta_t = O(1/\sqrt{t})$
- μ -Strongly convex: rate $O(1/(\mu t))$ for ASGD with $\eta_t = O(1/(\mu t))$

Both SGD and ASGD are **slow**. Best case is $O(1/t)$ convergence when f is strongly convex, while $O(e^{-\rho t})$ for FG

Recent results improve this:

- Bottou and LeCun (2005)
- Shalev-Shwartz et al (2007, 2009)
- Nesterov et al. (2008, 2009)
- Bach et al. (2011, 2012, 2014, 2015)
- T. Zhang et al. (2014, 2015)

- Gradient descent:

$$\theta^t \leftarrow \theta^{t-1} - \frac{\eta_t}{n} \sum_{i=1}^n \nabla f_i(\theta^{t-1})$$

$O(nd)$ iteration but linear convergence $O(e^{-\rho t})$ (strongly cvx case)

- Stochastic gradient descent:

$$\theta^t \leftarrow \theta^{t-1} - \eta_t \nabla f_{i_t}(\theta^{t-1})$$

$O(d)$ iteration but slow convergence $O(1/t)$ (strongly cvx case)

Do a fast algorithm with $O(d)$ iteration exist?

- Put $X = \nabla f_I(\theta)$ with I uniformly chosen at random in $\{1, \dots, n\}$
- We want to use Monte Carlo samples to approximate $\mathbb{E}X = \nabla f(\theta)$
- We find out C s.t. $\mathbb{E}C$ is easy to compute and such that C highly correlated with X
- Put $Z_\alpha = \alpha(X - C) + \mathbb{E}C$ for $\alpha \in [0, 1]$. We have

$$\mathbb{E}Z_\alpha = \alpha\mathbb{E}X + (1 - \alpha)\mathbb{E}C$$

and

$$\text{var } Z_\alpha = \alpha^2(\text{var } X + \text{var } C - 2\text{cov}(X, C))$$

- Standard variance reduction: $\alpha = 1$, so that $\mathbb{E}Z_\alpha = \mathbb{E}X$ (unbiased)

Idea: combine SGD with variance reduction

$$\theta^t \leftarrow \theta^{t-1} - \eta \left(\alpha (\nabla f_{i_t}(\theta^{t-1}) - \nabla f_{i_t}(\varphi^{t-1})) + \frac{1}{n} \sum_{i=1}^n \nabla f_i(\varphi^{t-1}) \right)$$

where $\nabla f_{i_t}(\varphi^{t-1})$ is the “last computed” gradient of ∇f_{i_t} along the iterations

- $\alpha = 1/n$: SAG (Stochastic Average Gradient, Bach et al. 2013)
- $\alpha = 1$: SVRG (Stochastic Variance Reduced Gradient, T. Zhang et al. 2015, 2015)
- $\alpha = 1$: SAGA (Bach et al., 2014)

Stochastic Average Gradient (SAG, Bach et al. 2013)

- **Input:** starting point θ_0 , learning rate $\eta > 0$
- For $t = 1, 2, \dots$ until *convergence* do
 - Pick at random (uniformly) i_t in $\{1, \dots, n\}$
 - Put

$$g_t(i) = \begin{cases} \nabla f_{i_t}(\theta^{t-1}) & \text{if } i = i_t \\ g_{t-1}(i) & \text{otherwise} \end{cases}$$

and compute

$$\theta^t = \theta^{t-1} - \frac{\eta}{n} \sum_{i=1}^n g_t(i)$$

- **Return** last θ^t

Variance reduction with SAG

Assume

- Each f_i is L -smooth
- f is μ -strongly convex
- $\eta_t = 1/(16L)$ constant
- Initialize using one epoch of SGD

Non-strongly convex case:

$$\mathbb{E}[f(\theta^t) - f(\theta_*)] \leq O\left(\frac{\sqrt{n}}{t}\right)$$

Strongly convex case:

$$\mathbb{E}[f(\theta^t) - f(\theta_*)] \leq O\left(\frac{1}{n\mu} + \frac{L}{n}\right) \exp\left(-t\left(\frac{1}{8n} \wedge \frac{\mu}{16L}\right)\right)$$

Improves a lot FG and SGD algorithms

- Complexity $O(d)$ instead of $O(nd)$ at each iteration
- Choice of a **fixed** step-size $\eta > 0$ possible
- But extra memory required: need to save all the previous gradients.

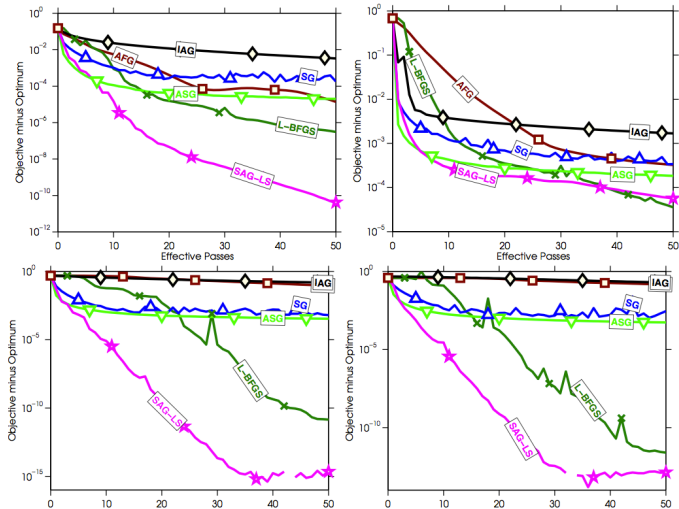
Hopefully

$$\nabla f_i(\theta) = \ell'(y_i, \langle x_i, \theta \rangle) x_i,$$

so only need to save $\ell(y_i, \langle x_i, \theta \rangle)$. Memory footprint is $O(n)$ instead of $O(nd)$. If $n = 10^7$, this is 76 Mo

Variance reduction with SAG

Comparison of convergence [Le Roux et al. 2012]



Now some practical problems / tricks around this

- Need to **index large data files** to be able to read lines at random fast

Many tools to do this. In Hadoop¹ there is the `ArrayFile.Reader` that does that. It's only (roughly) 3x slower than a sequential read of the file.

Stochastic optimization algorithm also work when using random shuffling of lines at beginning of each epoch: allows to further improve I/O.



Some practical remarks: lazy updates

Feature vectors are usually very sparse (words counts). Complexity of the iteration of a stochastic optimization algorithm can be reduced from $O(d)$ to $O(s)$, where s is the sparsity of the features. Important since $d \approx 10^6$ while $s \approx 10^3$

For minimizing

$$\frac{1}{n} \sum_{i=1}^n \ell(y_i, \langle \theta, x_i \rangle) + \frac{\lambda}{2} \|\theta\|_2^2$$

an iteration of SGD writes

$$\theta^t = (1 - \eta_t \lambda) \theta^{t-1} - \eta_t \ell'(y_i, \langle x_i, \theta^{t-1} \rangle) x_i$$

If x_i is s sparse, then computing $\eta_t \ell'(y_i, \langle x_i, \theta^{t-1} \rangle) x_i$ is $O(s)$, but $(1 - \eta_t \lambda) \theta^{t-1}$ is $O(d)$...

Some practical remarks: lazy updates

Trick: put $\theta^t = s_t \beta^t$, with $s_t \in [0, 1]$ and $s_t = (1 - \eta_t \lambda) s_{t-1}$

$$\theta^t = (1 - \eta_t \lambda) \theta^{t-1} - \eta_t \ell'(y_i, \langle x_i, \theta^{t-1} \rangle) x_i$$

becomes

$$\begin{aligned} s_t \beta^t &= (1 - \eta_t \lambda) s_{t-1} \beta^{t-1} - \eta_t \ell'(y_i, s_{t-1} \langle x_i, \beta^{t-1} \rangle) x_i \\ &= s_t \beta^{t-1} - \eta_t \ell'(y_i, s_{t-1} \langle x_i, \beta^{t-1} \rangle) x_i \end{aligned}$$

so the iteration is now

$$\beta^t = \beta^{t-1} - \frac{\eta_t}{s_t} \ell'(y_i, s_{t-1} \langle x_i, \beta^{t-1} \rangle) x_i$$

which has complexity $O(s)$.

Some practical remarks: lazy updates

- Just check that s_t is not too small once in a while, in this case put $s_t = 1$ and update θ_t and β_t
- Write the algorithm using only β_t , return $s_t\beta_t$ in the end

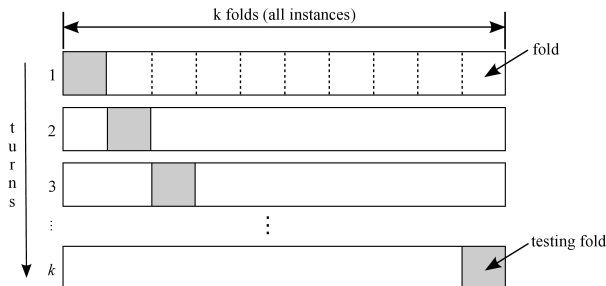
Now, complexity of one iteration of a stochastic algorithm is $O(s)$, while an approach based on FG methods is $O(nd)$ without using sparsity

Some practical remarks: cross-validation

Choice of penalization parameter λ by V -fold with SGD

Quick recap on V -fold:

- Take $V = 5$ or $V = 10$. Pick a random partition I_1, \dots, I_V of $\{1, \dots, n\}$, where $|I_v| \approx \frac{n}{V}$ for any $v = 1, \dots, V$



Some practical remarks: cross-validation

- I don't load the full data in memory
- I can't use V-Fold this way when I'm using an SGD-based solver

Simple solution: when picking a line i at random in the optimization loop, its fold number is given by $i \% V$

- Pick i uniformly at random in $\{1, \dots, n\}$
- Put $v = i \% V$
- For $v' = 1, \dots, V$ with $v' \neq v$: Update $\hat{\theta}_{v'}$ using line i
- Update the testing error of $\hat{\theta}_v$ using line i

Cross-validation: warm starting

- So I have many optimization problems to solve for choosing λ !
- If I'm using V -Fold cross-validation, and I choose a set $\Lambda = \{\lambda_1, \dots, \lambda_M\}$ of values for Λ , it is $V \times |\Lambda|$ problems
- But solutions $\hat{\theta}_{\lambda_{j-1}}$ and $\hat{\theta}_{\lambda_j}$ are going to be close when λ_{j-1} and λ_j are

Use **warm starts**:

- Fix parameters $\lambda_1 < \lambda_2 < \dots < \lambda_M$
- Put $\theta_0 = 0$ (I don't know where to start)
- For $m = M, \dots, 1$
 - Put $\lambda = \lambda_m$
 - Solve the problems starting at θ_0 for this value of λ (on each fold)
 - Keep the solutions $\hat{\theta}$ (test it, save it...)
 - Put $\theta_0 \leftarrow \hat{\theta}$

This allows to solve much more rapidly the sequence of problems