

MongoDB

Introduction

What is MongoDB

- Document-Oriented storage
 - The concept of a document replaces the row
- Utilizes “SQL” features
 - Index Support
- Easy to scale
 - Auto-Sharding
 - Auto-Balancing
- Querying
 - Native language : Javascript
- Map/Reduce

JSON

```
{  
  "interests": ["Exercise", "Family  
Time", "Skiing", "Camping"],  
  
  "name": "Barbara",  
  "dob": "10/01/1985",  
  "lastname": "Simmons",  
  "sex": "F",  
  "favorites": {  
    "color": "Red",  
    "sport": "Football",  
    "music": "HipHop/Rap"  
  }  
}
```

- Key-value representation in textual format
 - Human readable
 - Most programming languages translate it to dictionary-type data
 - In Javascript a json object works like a key-value dictionary
 - Order of keys not specific

Document store

RDBMS	MongoDB
Database	Database
Table, View	Collection
Row	Document (JSON,BSON)
Column	Field
Index	Index
Join	Embedded Document

Document:

```
{
  "_id" : ObjectId("5114e0bd42..."),
  "first" : "John",
  "last" : "Doe",
  "age" : 39,
  "interests" : [
    "Reading",
    "Mountain Biking ]
  "favorites": {
    "color": "Blue",
    "sport": "Soccer"}
}
```

Interacting With MongoDB

- Various APIs
 - Python, Ruby, Perl, Java, Java, Scala...
- MongoDB uses natively javascript
 - Simple language and easy to use
 - No datatypes
 - You can write Javascript in the shell or load a file to run
- UI: a lot off options
 - mongo-express, Edda, HumongouS, Umongo, MongoVision
- Here we are going to use (mostly)the shell
 - If you are writing a script in a file: field names should be in strings (single/double quoted)

Javascript

//variable declaration

```
var p=10;
```

//datatypes are decided on running time

```
p="text";
```

//while loop

```
var i=1;
```

```
while(i>0 && i<=10){
```

```
....
```

```
}
```

//if

```
if (condition) {
```

```
....
```

```
} else {
```

```
....
```

```
}
```

Logical operators :

- and: &&
- or: ||
- not : !

Javascript – Functions

```
function f(p1 , p2){  
    return p1*p2;  
}
```

```
var prod=f(5,6);  
print (prod.toString());  
//you need to convert a number  
//to string before you print it
```

```
//you can set a variable as a  
//function  
var f=function(p1 , p2){  
    return p1*p2;  
};
```

```
var prod=f(5,6);  
print (prod.toString());
```

Arrays

- Arrays can contain multiple types of data

```
text=""
var index;
var fruits = ["Banana", "Orange", "Apple", "Mango"];
for (index = 0; index < fruits.length; index++) {
    text += fruits[index];
}
```

```
text=""
var stuff = ["Banana", 5, [10,9]];
for (var index = 0; index < stuff.length; index++) {
    text += stuff[index].toString();
}
```


Starting MongoDB

- Standalone Server :

```
> mongod -port 2020
```

- You can define the port
- Other wise the default one is 27017

```
> mongo 127.0.0.1:2020
```

// OR

```
> mongo 127.0.0.1:2020/db
```

//if we already have a script e.g. test.js

```
> mongo 127.0.0.1:2020/db test.js
```

- The port has to be the same
- 127.0.0.1** : the local host/machine
- /db** : we can specify directly which database to use

Looping over JSON

```
var json={ "name":"John",  
  "age" : 19,  
  "rating": 1  
};
```

```
for (var key in json){  
    print (key+" "+json[key]);  
}
```

```
//you can print it fast and 'nice' with :  
printjson(json);
```

Loading the database

- If you started the command line :
 - Exit with Ctrl+d

Where the server is

Database to use (create it if it does not exist)

```
mongoimport --host localhost --port 6666 --db usersP --collection  
profiles --file ./data.json --jsonArray
```

Collection to use
(create it if it does not exist)

Path to json file

The file has multiple
documents into an Array

Insert

```
db.profiles.insert({name:"Rick",lastname:"Sanchez", dob:ISODate("1950-01-18T00:00:00Z")})
```

- You can insert stuff with completely different fields
 - A completely inconsistent “schema” would make the database management very hard

The cursor

```
var myCursor = db.profiles.find( { name: 'John' } );  
while (myCursor.hasNext()) {  
    printjson(myCursor.next());  
}
```

```
var myCursor = db. profiles.find( { name: 'John' } );  
  
myCursor.forEach(function (doc) { printjson(doc);});
```

Selection and Projection

- **db.collection.find(query, projection)**
 - Query and projection optional

```
db.profiles.find({name:{$eq:'John'}},{favorites:1})
```

```
db.profiles.find({name:{$eq:'John'}},{favorites:0})
```

```
db.profiles.find({}, {_id:0})
```

```
db.profiles.find({}, {_id:0}).limit(10)
```

```
db.profiles.find({$and:[{name:{$eq:'John'}},{\"favorites.color\":\"Black\"}]})
```

```
db.profiles.find({$and:[{name:{$eq:'John'}},{\"favorites.color\":{\"$ne:null}}]},  
{favorites:1})
```

Delete documents

- Delete on query

```
db.profiles.remove({name:"Rick",lastname:"Sanchez"})
```

Sort and Count

```
db.profiles.find({name:{$eq:'John'}},{name:1,lastname:1}).sort({lastname:1})
```

```
db.profiles.find({name:{$eq:'John'}},{favorites.color:1})  
                                .sort({'favorites.color':-1})
```

```
db.profiles.find({name:{$eq:'John'}}).count()
```


Update

- `db.collection.update(<query>, <update>, { upsert: <boolean>, multi: <boolean>,.})`

//insert him again

```
db.profiles.insert({name:"Rick",lastname:"Sanchez", dob:ISODate("1950-01-18T00:00:00Z")})
```

```
db.profiles.update({name:"Rick",lastname:"Sanchez"},{name:"Rick",lastname:"Sanchez",dob:ISODate("1951-01-18T00:00:00Z"),nephew:"Morty"})
```

- `$set/$unset`

//this does not require to re-write the entire document

```
db.profiles.update({name:"Rick",lastname:"Sanchez"},{$set:{occupation:"Scientist"}})
```

Arrays in MongoDB (1)

//We can set an array

```
db.profiles.update({name:"Rick",lastname:"Sanchez"},{$set:{interests:['bo  
oze','science']}})
```

//or add new stuff

```
db.profiles.update({name:"Rick",lastname:"Sanchez"},{$push:{interests:'m  
oney'}})
```

```
db.profiles.update({name:"Rick",lastname:"Sanchez"},{$push:{interests:['uni  
ty','rebellion']}})
```

```
db.profiles.update({name:"Rick",lastname:"Sanchez"},{$pushAll:{interests:['  
unity','rebellion']}})
```

Arrays in MongoDB (2)

We can query on the values of an array

e.g. How many people are interested in Reading ?

```
>db.profiles.find({interests:'Reading'}).count()
```

We can also search based on the length of the array

e.g. How many are interested in Reading AND have 2 or more interests?

```
>db.profiles.find({interests:'Reading', $where:"this.interests.length>=2"}).count()
```

How many are interested in Reading OR Walking?

```
db.profiles.find({interests:{$in:['Reading','Walking']}}).count()
```

How many are interested in Reading AND Walking?

```
db.profiles.find({interests:{$all:['Reading','Walking']}}).count()
```

Indexes

- By default there is an index on `_id`

```
//index in ascending order the values  
db.profiles.createIndex({name:1})
```

```
//allow only unique values  
db.profiles.createIndex({name:1}, { unique: true })
```

```
//build in background  
db.profiles.createIndex({name:1}, {background: true})
```

```
//compound  
db.profiles.createIndex({name:1,lastname:1})
```

Distinct values

- `db.collection.distinct(field, query)`

```
db.profiles.distinct('interests')
```

Aggregation in MongoDB

- `db.collection.aggregate([{ <stage> }, ...])`

```
db.profiles.aggregate([{$project:
                        {'age':
                          {$divide:[
                            {$subtract:[new Date(),'$dob']},
                            31558464000
                          ]
                        },
                        'name':1}
                      },
                      {$group:{
                        _id:'$name',
                        avgAge:{$avg:'$age'}
                      }
                      },
                      {$match:{avgAge:{$gt:18}
                      }}}])
```

Accumulators

\$sum	Returns a sum for each group. Ignores non-numeric values.
\$avg	Returns an average for each group. Ignores non-numeric values.
\$first	Returns a value from the first document for each group.
\$last	Returns a value from the last document for each group.
\$max	Returns the highest expression value for each group.
\$min	Returns the lowest expression value for each group.
\$push	Returns an array of expression values for each group.
\$addToSet	Returns an array of unique expression values for each group. Order of the array elements is undefined.

Operator (reminder)

Name	Description
\$project	Manage the fields you want to use
\$match	Apply a query to filter the data
\$limit	Use only the first n documents
\$skip	Skip n documents
\$unwind	Applied in a array which flattens it .
\$group	Groups input documents by a specified identifier expression and applies accumulator expression(s),.
\$sort	Reorders the document stream by a specified sort key.
\$out	Writes the resulting documents of the aggregation pipeline to a collection.

Aggregation in MongoDB

Find the top 5 most popular interests :

```
db.profiles.aggregate(  
[  
  {$unwind: "$interests"},  
  { $group:  
    { _id: "$interests",  
      "count": { $sum: 1 } }  
  },  
  {$sort :{"count":-1}},  
  {$limit: 5}  
])
```

If the aggregating field is an array we can expand so all possible values are used

Each appearance of a keyword counts as 1
Sum on the appearances

Sort in descending order by appearance

Keep only the first 5 results

Exercises

1. Count how many people are over 30
2. In the same aggregation:
 - For each music genre count the number of people who favorite it and group music genres in bin-counts of '10s'
 - E.g. Pop and rock were like by [10-20) then they belong to bin 10
 - Helpful functions: $\$mod[a,b] = \text{modulo}(a,b)$, $\$subtract(a,b)=a-b$
 - $\text{floor}(n)=n - n \bmod 1$
3. Find the top-5 music genres for people who like "Listening to Music"
4. For each favorite sport find the top interest

Resources

- <http://docs.mongodb.org/manual/>
- Python (almost exactly the same API as javascript):
 - <https://api.mongodb.org/python/current/>
- Java :
 - <http://docs.mongodb.org/getting-started/java/>
- UIs:
 - <http://docs.mongodb.org/ecosystem/tools/administration-interfaces/>