



# Parallel Processing using MapReduce

Prof. Yanlei Diao

Some materials borrowed from talks by Jeff Dean, Sanjay Ghemawat, Christophe Bisciglia, Aaron Kimball, and Sierra Michels-Slettvet



# Motivation: Large Scale Data Processing

---

- Want to process lots of data, *unstructured or structured*
- Want to parallelize across *hundreds/thousands* of commodity computers
  - New definition of cluster computing: *large numbers of low-end processors working in parallel to solve a computing problem.*
  - Parallel DB: *a small number of high-end servers.*
- Want to make this easy

•  
•  
•  
•  
•  
•  
•

# I. MapReduce

---

- Clean abstraction for programmers
- Automatic parallelization & distribution
- Fault-tolerance
- Status and monitoring tools

MapReduce: Simplified Data Processing on Large Clusters. Jeffrey Dean and Sanjay Ghemawat. OSDI 2004.

# Programming Model

---

- Borrows from functional programming
- Users implement an interface of two functions:
  - `map (in_key, in_value) -> list(out_key, intermediate_value)`
  - `reduce (out_key, list(intermediate_value) -> list(out_value)`

# map

---

- Input: a key-value pair. E.g.,
  - A line out of files (filename, line),
  - A row of a database (row\_id, row),
  - A document (doc\_name, document)
- map( ) produces one or more *intermediate* values along with an output key from the input.
- map( ) is **stateless**: one input leaves no state that would affect the processing of the next input.

# reduce

---

- After the map phase is over, **all** the intermediate values for a given output key are collected into a list
- `reduce( )` combines those intermediate values into one or more *final values* for that same output key
- `reduce( )` can be **stateful**: it operates on all the intermediate values of a certain key

# Example: Count Word Occurrences

---

```
map(String input_key, String input_value):
```

```
    // input_key: document name
```

```
    // input_value: document contents
```

```
    for each word w in input_value:
```

```
        EmitIntermediate(w, "1");
```

```
reduce(String output_key, Iterator intermediate_values):
```

```
    // output_key: a word
```

```
    // output_values: a list of counts
```

```
    int result = 0;
```

```
    for each v in intermediate_values:
```

```
        result += ParseInt(v);
```

```
    Emit(AsString(result));
```

*How do we implement this using a relational DBMS? Customized data loading (data may be used only once), then Group By.*

# Click Stream Analysis: Page Frequencies

---

Clicks(time, url, referral\_url, user\_id, geo\_info...)

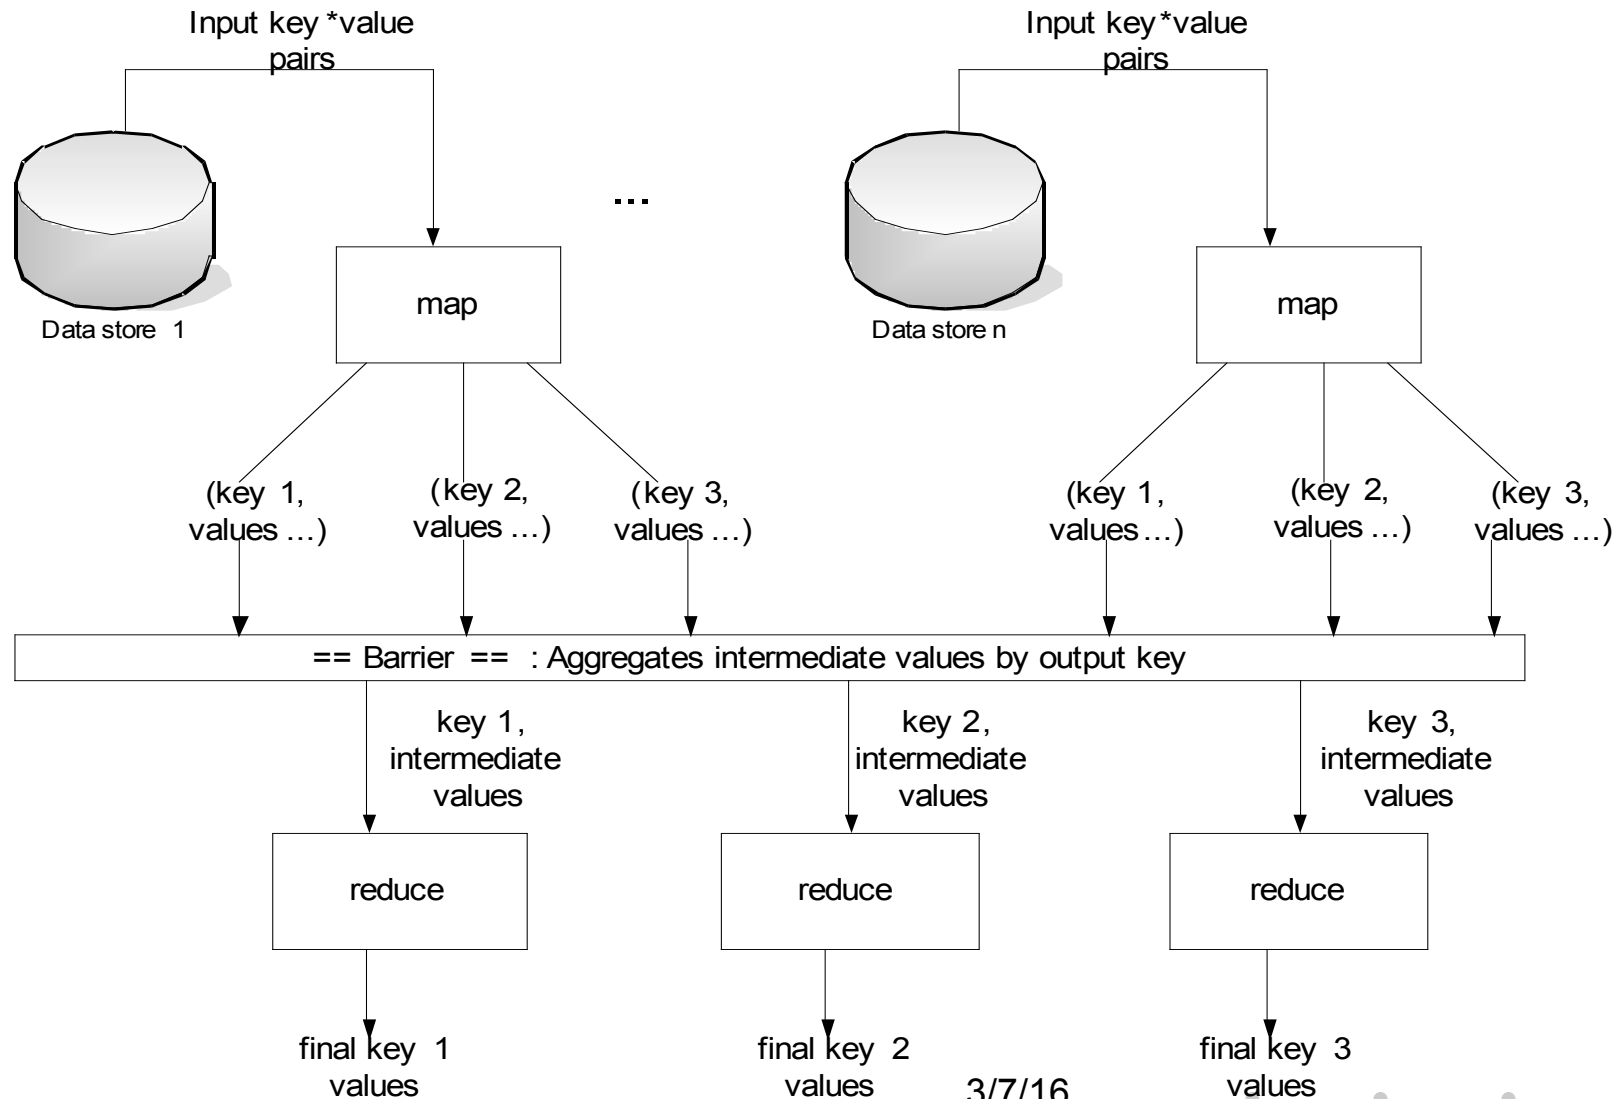
```
map(String tuple_id, String tuple):  
    EmitIntermediate(url, "1");  
  
reduce(String url, Iterator list_tuples):  
    int result = 0;  
    for each t in list_tuples:  
        result += ParseInt(t);  
    Emit(AsString(result));
```

```
Select count(*)  
From Clicks  
Group By url;
```



# MapReduce Computation Model

Extract (key, value) using map(). **Group data by key.** Then apply reduce().



# Parallelism

---

- The map() function is stateless, so many instances can run **in parallel on different splits** (chunks) of input data
- The reduce() function is stateful, but works on an output key at a time, so many copies can run **in parallel on different keys (groups)**
- Performance bottleneck: reduce phase can't start until map phase is completely finished.

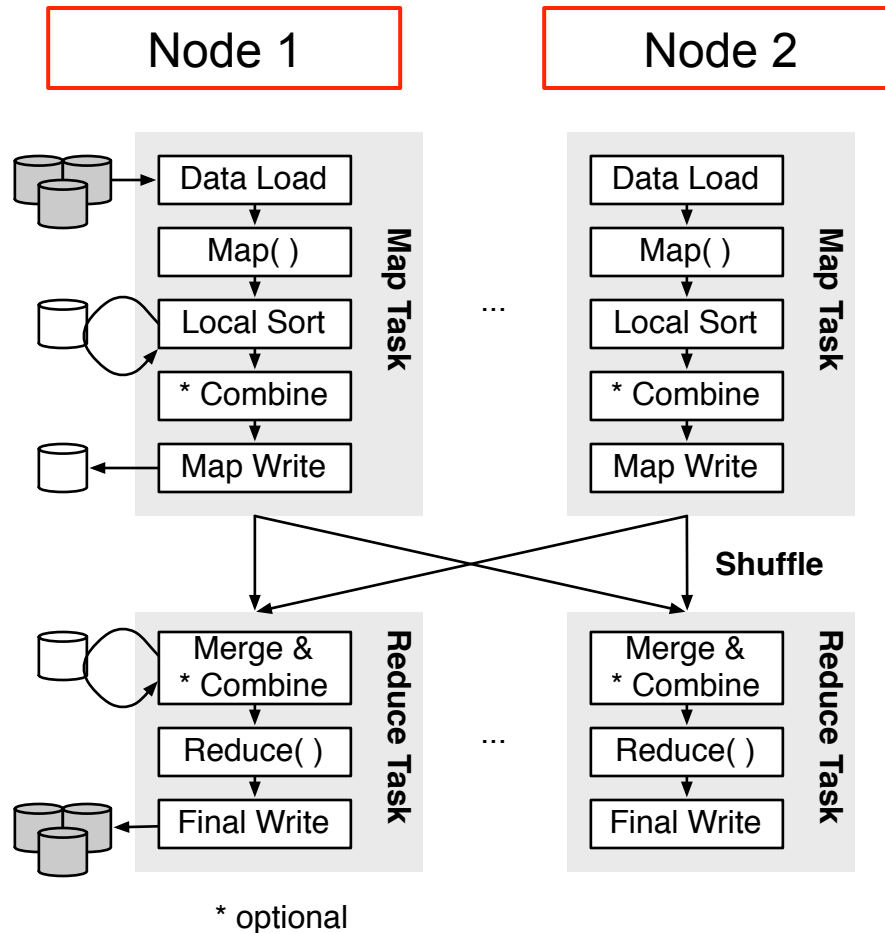
# Optimization: Incremental Computation

---

- “Combiner” functions can be applied earlier, e.g., right after map() finishes on the same machine
- Causes a mini-reduce phase to occur before the real reduce phase, to save bandwidth
- Common examples: word frequency, url frequency
- Also called partial aggregation

*Under what conditions is it sound to use a combiner?*

# Illustration of the Dataflow



# Fault Tolerance

---

- Fine-grained fault tolerance: **materialize** map output onto local disk before the map task completes
- Master detects worker failures
  - Re-executes completed & in-progress **map()** tasks
  - Re-executes in-progress **reduce()** tasks
- Master notices particular input key/values cause crashes in **map()**, and skips those values on re-execution.
  - Effect: Can work around bugs in third-party libraries!

# Optimization: Redundant Execution

---

- No reduce can start until map is complete:
  - A single slow disk controller can rate-limit the whole process
- Master redundantly executes “slow-moving” map tasks; uses results of first copy to finish

# Refinement: Exploiting Locality

---

- Master scheduling policy:
  - Asks GFS for locations of replicas of input file blocks
  - Map tasks typically split into 64MB (GFS block size)
  - Map tasks scheduled so GFS input block replica are on same machine or same rack
- Effect
  - Thousands of machines read input at local disk speed
  - Without this, rack switches limit read rate

## II. Comparison to Parallel Databases

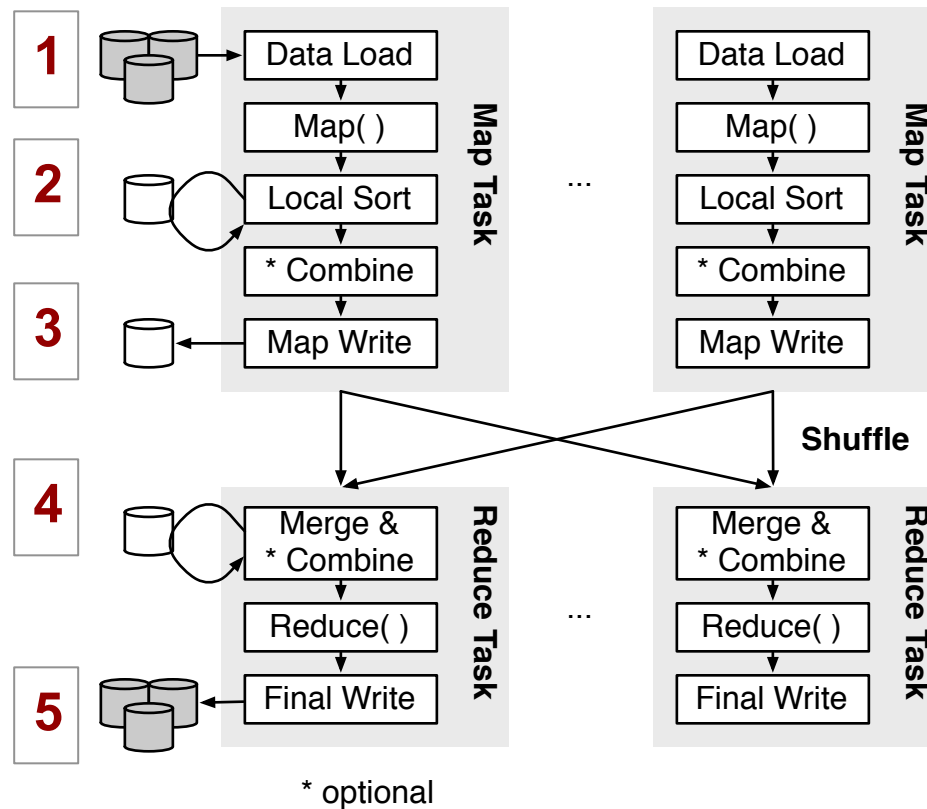
---

- Let us consider structured data here.
  - Of course, MapReduce can also handle text processing!
- 1. A closer look at internal implementation of MapReduce
  - Extract (key, value) using map()
  - **Group data by key**
  - Then apply reduce() to each group
- 2. Implementing relational operators using MapReduce
  - Parallel sorting?
  - Parallel Join?
  - Parallel group by-aggregation?
- 3. MapReduce query plans



# 1. Analysis of Open Source Hadoop

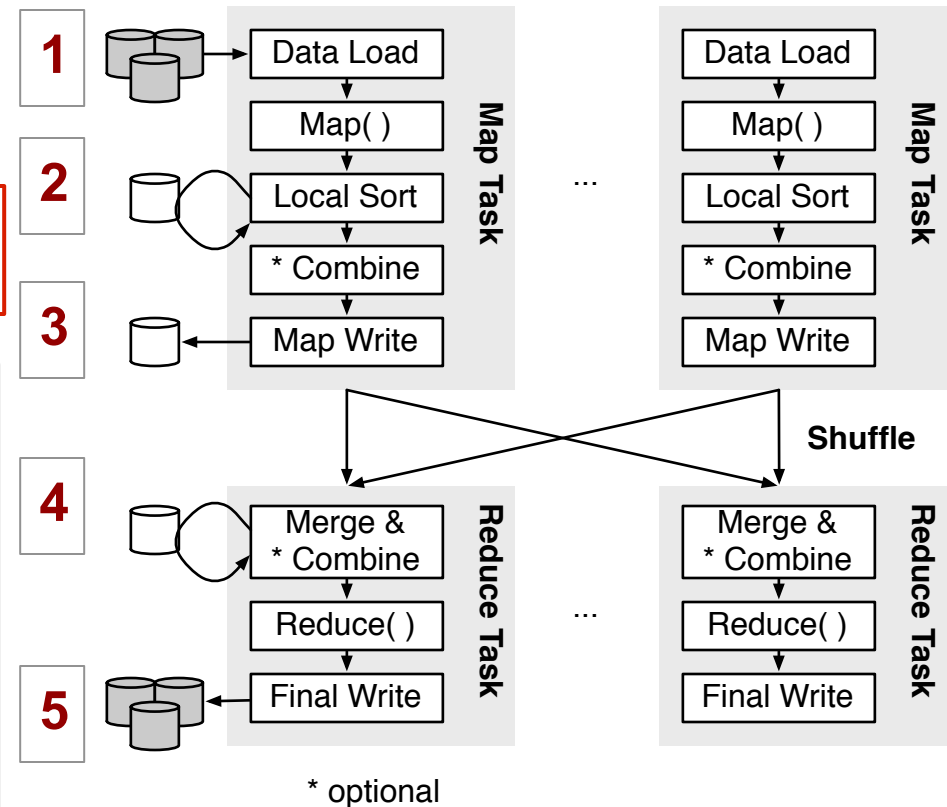
## Sort-Merge Implementation of Group-By



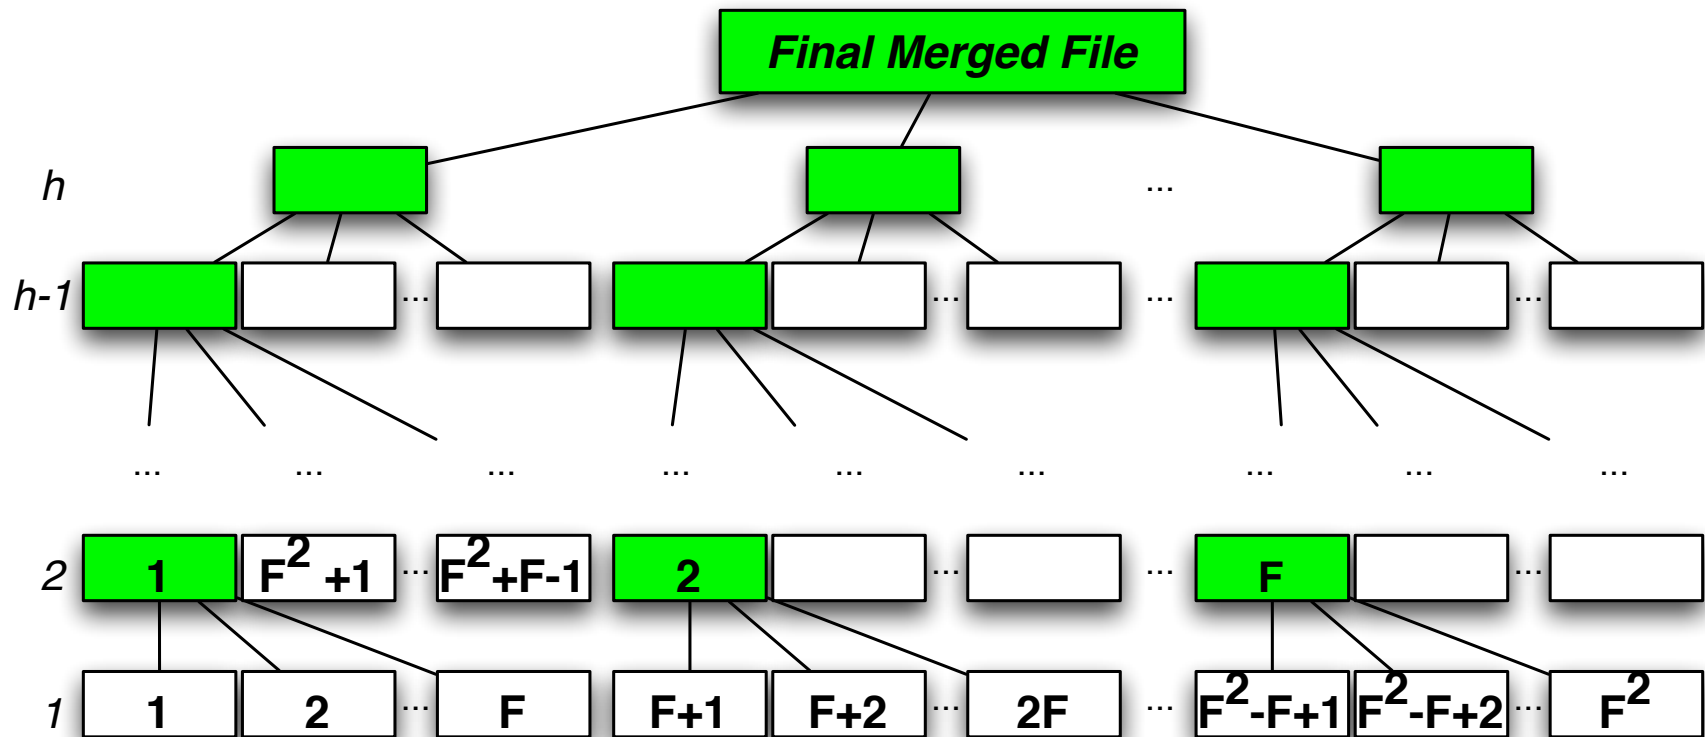
# Analysis of Open Source Hadoop

|          |   |
|----------|---|
| <b>R</b> | Number of reduce tasks per node                               |
| <b>C</b> | Map input chunk size  |
| <b>F</b> | Merge factor that controls how often on disk files are merged |

| Symbol                                  | Description   |
|---|---|
| <b>(1) System Settings</b>              |   |
| $R$                                     | Number of reduce tasks per node   |
| $C$                                     | Map input chunk size  |
| $F$                                     | Merge factor that controls how often on-disk files are merged   |
| <b>(2) Workload Description</b>         |   |
| $D$                                     | Input data size   |
| $K_m$                                   | Ratio of output size to input size for the map function   |
| $K_r$                                   | Ratio of output size to input size for the reduce function  |
| <b>(3) Hardware Resources</b>           |   |
| $N$                                     | Number of nodes in the cluster  |
| $B_m$                                   | Output buffer size per map task   |
| $B_r$                                   | Shuffle buffer size per reduce task   |
| <b>(4) Symbols Used in the Analysis</b> |   |
| $U$                                     | Bytes read and written per node, $U = U_1 + \dots + U_5$ where $U_i$ is the number of bytes of the following types<br>1: map input; 2: map internal spills; 3: map output;<br>4: reduce internal spills; 5: reduce output |
| $S_i$                                   | Number of sequential I/O requests per node for IO type $i$  |
| $T$                                     | Time measurement for startup and I/O cost   |
| $h$                                     | Height of the tree structure for multi-pass merge   |

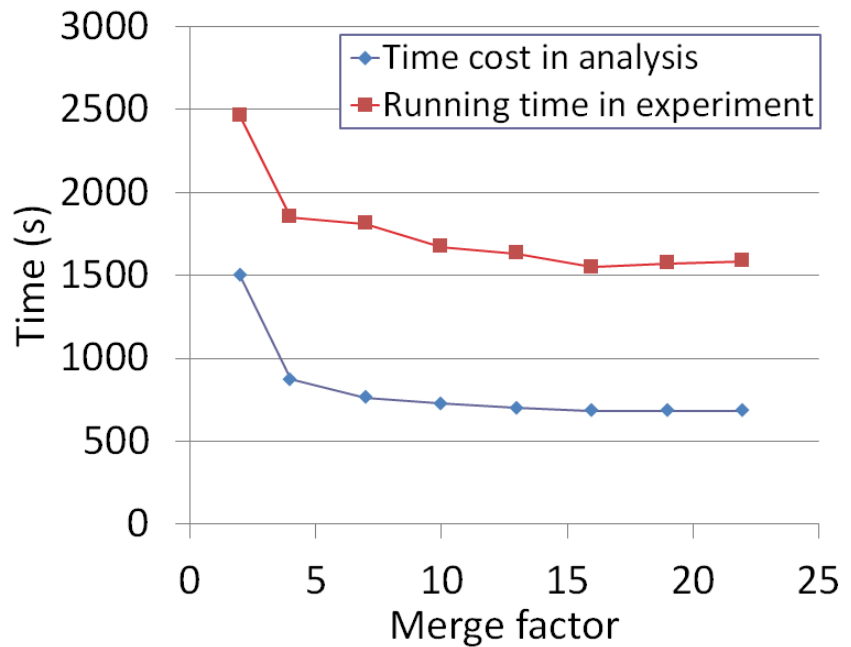
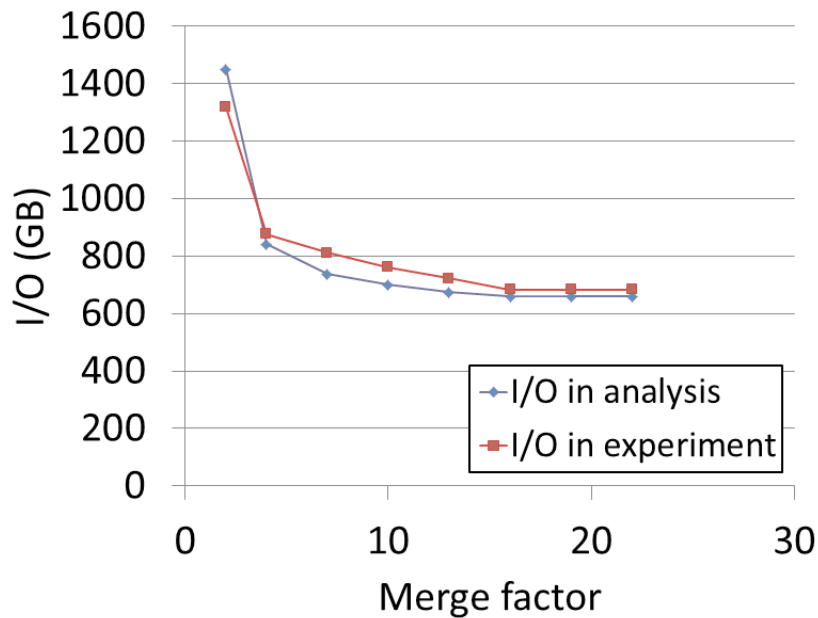


# Analysis of Multi-Pass Merge



- Used in a mapper for sorting if map output exceeds memory size
- Used in a reducer unless all data fits in memory

# Effect of the Merge Factor F



## 2. Implementing Relational Operators

---

- Selection:  $R.a > \text{"abc"}$ 
    - **ParallelDB**: if range partitioned, use a few nodes and indexes
    - **MapReduce**: scan all nodes, map() only.
      - *Can be dominated by start-up cost. No indexes in the original impl.*
  - Most other operators need *repartitioning* data:
    - **ParallelDB**: explicit partitioning function
    - **MapReduce**: more complicated
      - (1) Implicit partitioning function, **fn**, controls data shuffling to reducers.  
(Default is hash partitioning. Can be changed to range partitioning.)
      - (2) Each reducer uses an additional mechanism to group data by the key.
- ❑ Consider the task to range partition data and sort data in each range.  
What is the key in the MR programming model?

# Join Operators

---

- Equijoin:  $R.a = S.a$ 
  - **ParallelDB**: hybrid hash join.
    - I/O and network costs?
  - **MapReduce**: the programming interface is not natural for joins.
    1. `map()` annotates tuples with 'r' and 's',
    2. the system groups all data by the join attribute using sort-merge,
    3. `reduce()` joins 'r' and 's' tuples with the same value of the join attribute.
    - *It is better to change the programming model to make join more natural!*
- Non equijoin:  $R.a < S.a$ 
  - **ParallelDB**: fragment-replication
  - **MapReduce**: simulates fragment-replication. If replicate S,
    - replicate each S tuple  $m$  times in the mapper
    - tweak the partitioning function, **fn**, for shuffling so that these  $m$  copies go to different reducers (fn can be customized in Hadoop)

# Group By Aggregation

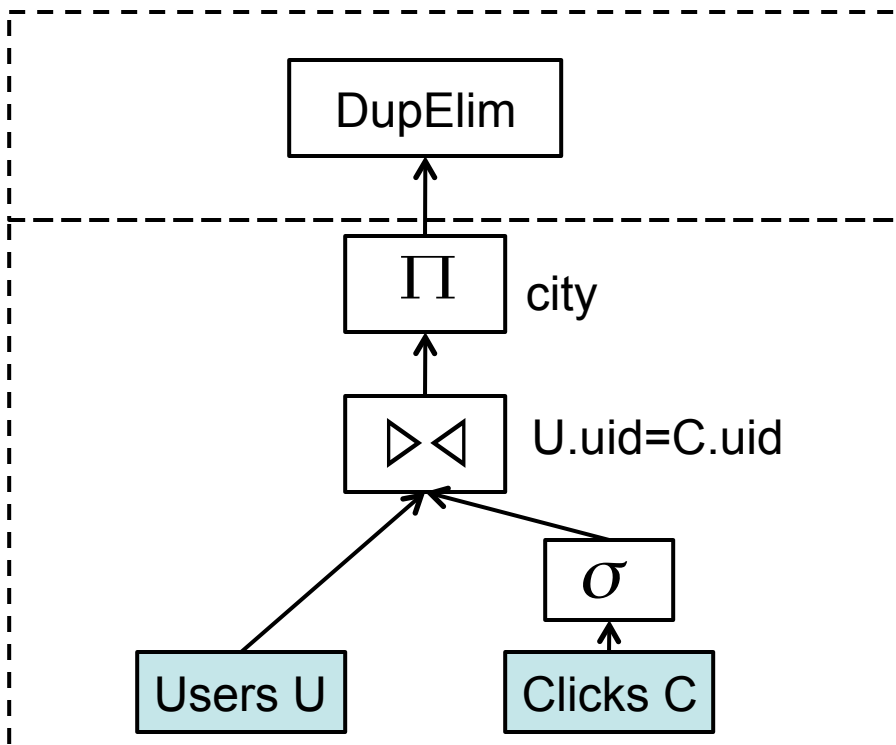
---

- Scalar aggregate: `count()`, `sum()`
  - **ParallelDB**: partial aggregation + final aggregation
  - **MapReduce**: `map()` is empty; use `combiner()` for partial aggregation; use `reduce()` for final aggregation
- Group by aggregation:  $G_{R.a, \text{aggr}(R.b)}$ 
  - **ParallelDB**: unary input version of hybrid hash join
  - **MapReduce**:
    - `map()` simply emits tuples;
    - the system groups data by `R.a`;
    - `reduce()` computes sum.
    - should use the `combiner()` for partial aggregation earlier.

### 3. MapReduce Query Plans

- How many rounds of map reduce jobs?
- In each round, what is in map(), what is in reduce()?

```
SELECT DISTINCT U.city  
FROM   Users U, Clicks C  
WHERE  U.uid=C.uid  
       AND C.url LIKE '%google%';
```



#### Round 2:

**Key:** city

**Map:** emit

**Reduce:** emit a tuple in each group

#### Round 1:

**Key:** uid

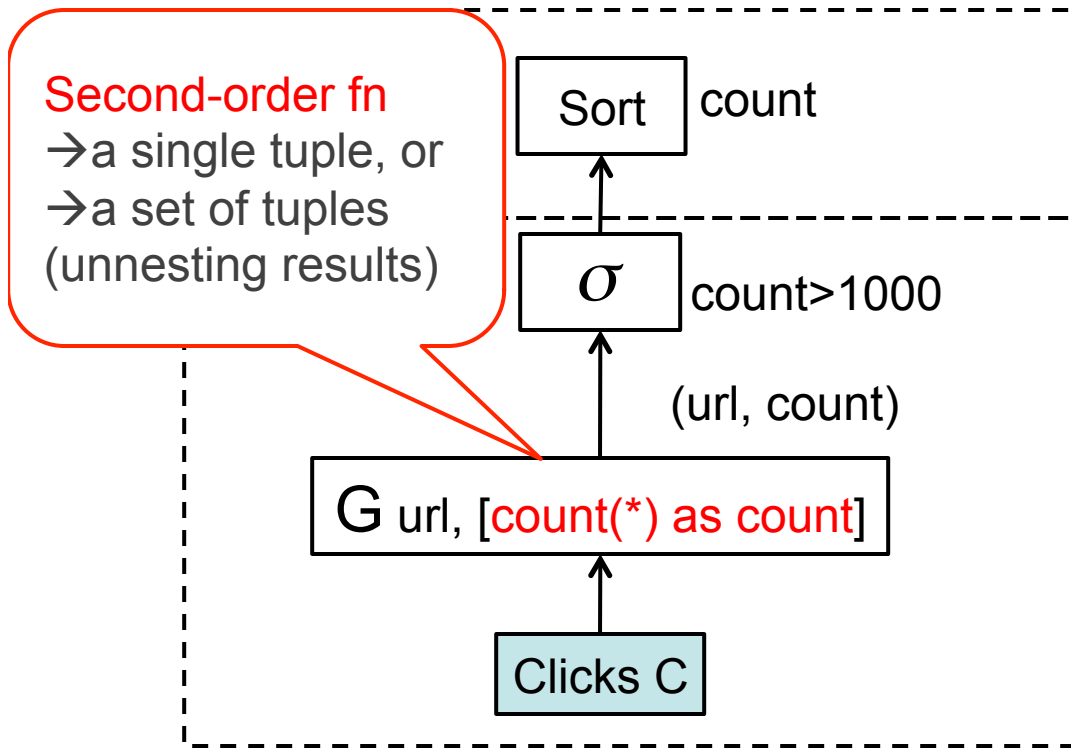
**Map:** (1) selection, (2) create 'u', 'c' tuples with labels

**Reduce:** (1) join tuples within each group, (2) emit cities



# More on Query Plans

```
SELECT url, count(*)
FROM Clicks C
GROUP BY url
HAVING count(*) > 1000
ORDER BY count(*) DESC;
```



## Round 2:

**Key:** count

**Map:** emit

**Shuffle:** range partitioning (set manually)

**Reduce:** local sort  
or (a simple but bad plan)

( **Key:** fixed

**Map:** emit

**Reduce:** sort all in a single reducer)

## Round 1:

**Key:** url

**Map:** emit

**Reduce:** (1) count, (2) selection, (3)  
emit (url, count)

# References

---

- MapReduce: Simplified Data Processing on Large Clusters. Jeffrey Dean and Sanjay Ghemawat. OSDI 2004.
- MapReduce and Parallel DBMSs: Friends or Foes?. Michael Stonebraker, Daniel Abadi, David J. DeWitt, Sam Madden, Erik Paulson, Andrew Pavlo, Alexander Rasin. CACM Jan 2010.
- MapReduce: A Flexible Data Processing Tool. Jeffrey Dean, Sanjay Ghemawat. CACM Jan 2010.
- Dryad/Linq: <http://research.microsoft.com/en-us/projects/dryadlinq/>
- Hyracks: A Flexible and Extensible Foundation for Data-Intensive Computing. Vinayak Borkar, et al. ICDE 2011.
- Llama: Leveraging Columnar Storage for Scalable Join Processing in the MapReduce Framework. Yuting Lin, etc. SIGMOD 2011.
- A Platform for Scalable One-Pass Analytics. Boduo Li, et al. SIGMOD 2011.

- 
- 
- 
- 
- 
- 
- 

# Questions

---

