



# Parallel Databases

Prof. Yanlei Diao



# A Data Warehouse for A Social Network

---

**User profiles:**  
100 Million users  
Each with profile,  
pics, postings,...

Web  
Server

Web  
Server

...

Web  
Server

**Click Streams:**  
1 billion rows/day  
5-10 TB/day

**Data Loading:**  
High Volume +  
Transformation

Data Processing  
Backend

**Quick lookups and updates:**  
Update your own profile, read  
friends' profiles, write msgs,...

**Analysis Queries:**  
Ad targeting, fraud detection,  
resource provisioning...

# Fun (Old) Numbers about Facebook

---

500 million active users

9.5% Internet traffic

>30,000 servers

Initial software:  
PHP + MySQL cluster  
+ Memcached



One of the largest  
MySQL cluster

>4.5 billion msgs/day  
>15 TB click logs/day

Stores >20 billion  
photos, and serves 1  
million img/sec.

# Parallel Databases 101


---

- Rise of parallel databases: late 80' s
- Architecture: shared-nothing systems
  - A number of nodes connected by fast Ethernet switches
  - But used special-purpose hardware (costly, slow to evolve)
  - Small scale (hence did not focus on fault tolerance)
- Typical systems
  - Gamma: U. of Wisconsin Madison
  - TeraData: Wal-Mart's 7.5TB sales data in hundreds of machines
  - Tandem
  - IBM / DB2
  - Informix...

# Some Parallel ( | | ) Terminology


- Speed-Up

- Holds the problem size, grows the system
- Reports  $\text{serial time} / \text{||-time}$
- Ideally, linear



- Scale-Up


- Grows both the system and the problem, reports running time
- Ideally, constant



# Different Types of DBMS ||-ism

---


- **Partitioned (data) parallelism**
  - Partition data over all nodes, get the nodes working to compute a given operation (scan, sort, join)
- **Pipelined parallelism**
  - A chain of operators  $O_1, O_2, \dots, O_k$  run in parallel, with  $O_1$  working on tuple  $t_n$ ,  $O_2$  on  $t_{(n-1)}$ , ...  $O_k$  on  $t_{(n-k+1)}$
  - Can run these operators on different nodes
  - Some operators break pipelining, e.g. sort, hash
- **Independent operators**
  - Consider bushy query plans
  - A join B, C join D are independent
- We'll focus on partitioned ||-ism



# Data Partitioning Schemes

Partitioning a table:


## Range



Good for seq. scan,  
associative search, sorting

Can have data skew


## Hash



Good for seq. scan, equality  
search, equijoins if they  
match the hash attr

Bad for range search, or  
operations that do not  
match the hash attr;  
Can also have data skew

## Round Robin



Good for seq. scan

Useless for other  
query operations

# Parallel Scans and Associative Accesses


---

- Scan in parallel, and merge.
- Selection may not require all sites for range or hash partitioning.
  - Want to restrict selection to a few nodes, or **restrict “small” queries to a few nodes.**
  - Indexes can be built at each partition.
  - What happens during data inserts and lookups?



# Parallel Sorting

---




- Some record in the history:
  - 8.5 Gb/minute, shared-nothing; Datamation benchmark in 2.41 secs (UCB students! <http://now.cs.berkeley.edu/NowSort/>)
- Idea:
  - Scan in parallel, and range-partition as you go.
  - As tuples come into each node, begin “local” sorting
  - Resulting data is sorted, and range-partitioned.
  - Problem: *skew!*
  - Solution: “sample” the data at start to determine partition points.

# Partitioned Join

---


- For equi-joins, *partition* the two input relations across all nodes, and compute the join locally at each processor.
- Can use either *range partitioning* or *hash partitioning*, on the join attribute
  - $R$  and  $S$  each are partitioned into  $n$  partitions, denoted  $R_0, R_1, \dots, R_{n-1}$  and  $S_0, S_1, \dots, S_{n-1}$ .
  - Partitions  $R_i$  and  $S_i$  are sent to node  $i$ .
  - Each node locally computes the join using any method.

# Parallel Hash Join



- In first phase, partitions get distributed to different nodes:
  - A good hash function *automatically* distributes work evenly!
- Do second phase at each node.
- Almost always the winner for equi-join.

# Dataflow Network for Parallel Join




- Good use of split/merge makes it easier to build parallel versions of sequential join code.

# Parallel Hybrid Hash Join

---

- Run the analysis of hybrid hash join as before, but using the aggregate memory  $B^*$ , to determine the number of *logical* buckets  $N$ ; hash fn  $h_1$  maps tuples to buckets.
- The in-memory join,  $R_0 \bowtie S_0$ , is partitioned over  $M$  processors using hash fn  $h_2$ .
- Disk resident buckets,  $R_i$  and  $S_i$  ( $i > 0$ ), are partitioned over  $K$  disks using  $h_3$ .
- For each subsequent bucket  $i$  ( $i > 0$ ), partition  $R_i \bowtie S_i$  over  $M$  processors using  $h_2$ .

# Parallel Hybrid Hash Join

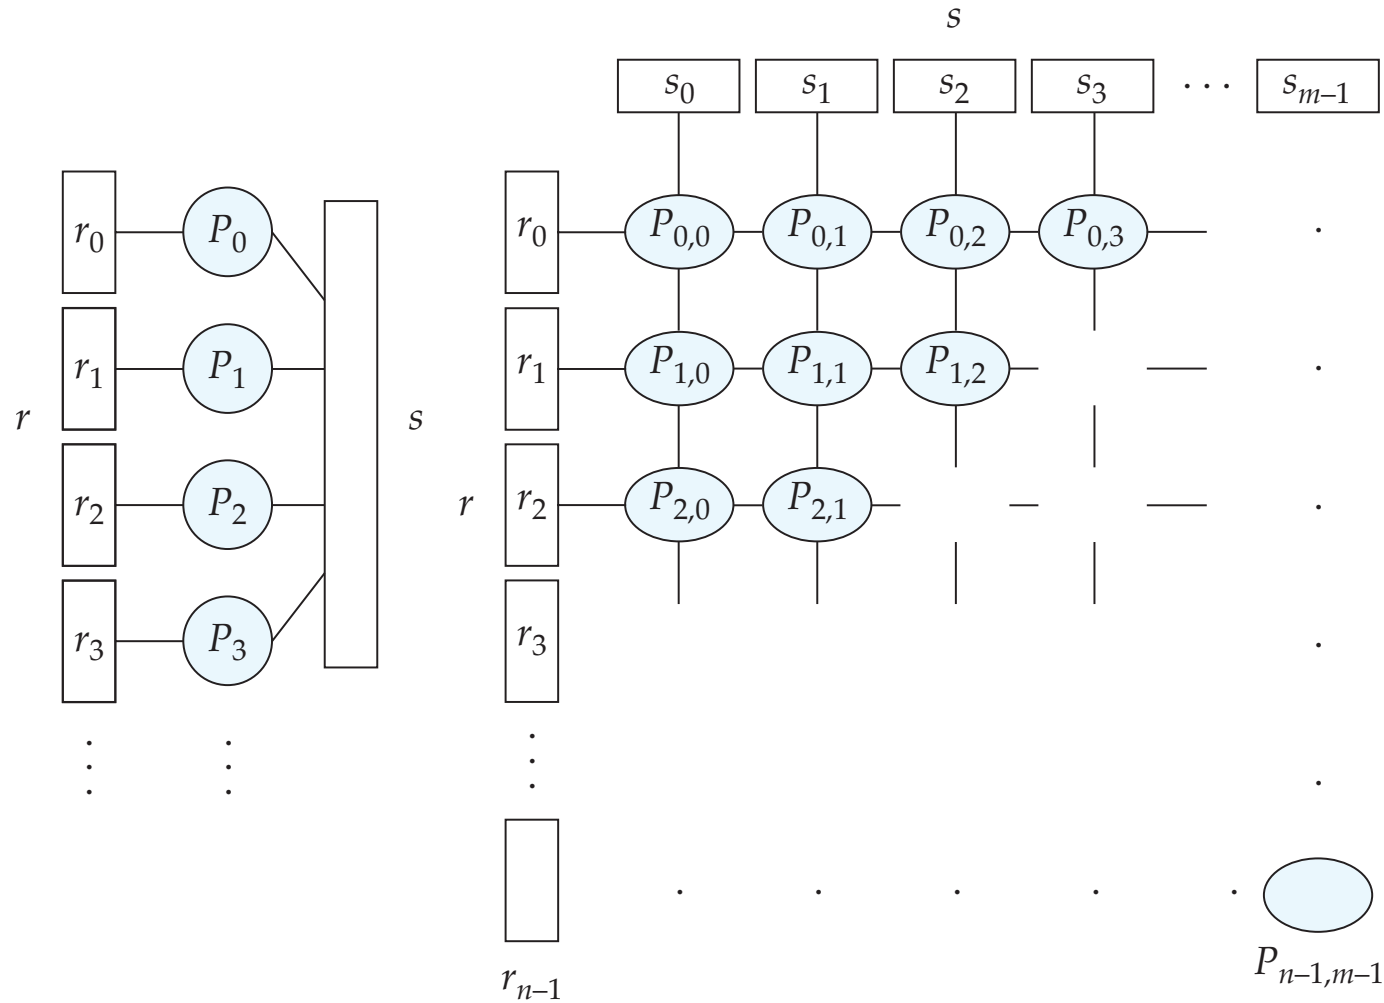


# Fragment-and-Replicate Join

---

- Partitioning not possible for some join conditions
  - E.g., non-equijoin conditions, such as  $R.A > S.B$ .
- Use the **fragment and replicate** technique
  - 1) Special case: only one relation is partitioned
    - $R$  is partitioned; any partitioning technique can be used.
    - The other relation,  $S$ , is replicated across all the nodes.
    - Node  $i$  then locally computes the join of  $R_i$  with all of  $S$  using any join technique.
    - Works well when  $S$  is small.
  - 2) General case: both  $R$  and  $S$  are partitioned
    - Need to replicate all  $R$  partitions or all  $S$  partitions
    - Depicted on the next slide

# Illustrating Fragment-and-Replicate Join



(a) Asymmetric  
fragment and replicate

(b) Fragment and replicate



# Parallel Aggregates

---

- For each aggregate function, need a decomposition:
  - **Distributive:**  $\text{count}(S) = \Sigma \text{count}(s(i))$ , ditto for  $\text{sum}()$
  - **Algebraic:**  $\text{avg}(S) = (\Sigma \text{sum}(s(i))) / \Sigma \text{count}(s(i))$
  - **Holistic:** e.g., median, quantiles
- For group-by aggregation:
  - How would you implement parallel group by?
  - How do you add aggregation to each group?

- 
- 
- 
- Questions
- 
- 
- 

---

