

-
-
-
-
-
-
-

Foundation of Relational Databases

Yanlei Diao

Slides Courtesy of R. Ramakrishnan and J. Gehrke



Teaching Staff

❖ Instructor: Yanlei Diao

- Email: yanlei.diao@polytechnique.edu
- How to address in email: “Dear Prof. Diao”, “Dear Yanlei”

❖ Teaching Assistants

- Michaël Thomazo michael.thomazo@inria.fr
- Oscar santiago Mendoza rivera
oscar-santiago.mendoza-rivera@inria.fr
- Prosper burq prosper.burq@gmail.com

Course Web Site

<http://datascience-x-master-paris-saclay.fr/>

Follow the link to **Moodle:**

- post lecture notes
- collect submissions

Textbook



Database Management Systems 3rd Edition

Ramakrishnan and Gehrke

Amazon:

- Buy new: \$43-\$147.09 (hardcover); paperback, \$23;
Kindle, rent options are also available...

Lecture notes will be posted in **Moodle**
after class.

Academic Honesty

- ❖ All submitted work must be your own!
 - Although students are encouraged to study together, each student must produce his or her own solution to each homework.
 - Copying or using **sections of someone else's program or assignment** (even if it has been modified by you), or copying a **solution from an external source**, is not acceptable.
 - The teaching staff will be vigorous in enforcing them.

Databases and DBMS' s

- ❖ A **database** is a large, integrated collection of data
- ❖ A **database management system** (DBMS) is a software system designed to store and manage a large amount of data
 - **Declarative interface** to define data stored, add data, update data, and query data
 - **Efficient querying**
 - **Concurrent users**
 - **Reliable storage** and **crash recovery**
 - **Access control...**

Early DBMS' s

- ❖ Early DBMS' s (1960' s) evolved from file systems
- ❖ Many small data items, many queries and updates
 - Banking, Airline reservations
- ❖ 1960s Navigational DBMS
 - Tree-based or graph-based data model
 - Manual navigation to find what you want
 - No support for “search” (≠ “program”)
- ❖ 1973 Turing Award Winner
 - Charles William Bachman
 - “The Programmer as Navigator”
 - The network data model




Relational DBMS

❖ Relational model (1970)

- **Data independence**: hides details of physical storage from users
- **Declarative query language**: say **what** you want, not **how** to compute it
- **Mathematical foundation**: what queries mean, possible implementations

❖ 1981 Turing Award Winner



- Edgar F. (“Ted”) Codd
- Mathematically-inclined researcher
- Legitimized DBMSs as a theoretically respectable research field in CS



Relational DBMS



1974 Debate at an
ACM Workshop



- ❖ Query optimization (1970' s till now)
 - Earliest: System R at IBM, INGRES at UC Berkeley
 - Queries can be **efficiently executed** despite data independence and declarative queries!
- ❖ 2014 Turing Award Winner
 - Michael Stonebraker
 - “For fundamental contributions to modern database systems”



Commercial DBMS' s

INGRES

UC Berkeley,
Stonebraker et al

Informix

Postgres

Sybase

MS SQL Server

System R

IBM San Jose,
Gray, Selinger et al

IBM DB2

Oracle

MySQL

Banking

Ticketing

Data
warehouse

E-commerce

Social net



Foundation of Relational Databases

Foundation of Relational Databases

- ❖ **Relational Model**
- ❖ Formal Query Languages
 - Relational Algebra
 - Relational Calculus
 - Language Theory

Relational Model

- ❖ A relational database is a set of *relations*.
- ❖ Each relation has:
 - *Schema* : specifies name of relation, name and type (domain) of each attribute.
 - Students(*sid*:string, *name*:string, *login*:string, *age*:integer, *gpa*:real).
 - *Instance* : a table with rows (*tuples*) and columns (*attributes, fields*).
cardinality = #rows, *degree / arity* = #columns.
- ❖ A relation is a *set* of tuples (in theory).
 - All rows must be distinct, no duplicates.

Example Instance of Students Relation

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

- ❖ Cardinality = 3, degree = 5
- ❖ All rows are distinct.
- ❖ Some columns of two rows can be the same.

Creating Relations in SQL

- ❖ Create the Students relation
- ❖ Specify *domain constraints*:
 - type of each field
 - later enforced by the DBMS upon tuple insertion or update.

```
CREATE TABLE Students
    (sid    CHAR(20),
     name   CHAR(20),
     login  CHAR(10),
     age    INTEGER,
     gpa    REAL);
```

```
CREATE TABLE Enrolled
    (sid    CHAR(20),
     cid    CHAR(20),
     grade  CHAR(2));
```

Adding Tuples

❖ Can insert a single tuple using:

```
INSERT INTO Students (sid, name, login, age, gpa)  
VALUES ( '53688' , 'Smith' , 'smith@ee' , 18, 3.2);
```

➡ *Powerful variants of these commands are available; more later!*

Integrity Constraints

- ❖ Integrity Constraints (IC's): condition that must be true for **any** instance of the database.
 - *Domain constraint*
 - *Primary key constraint*
 - *Foreign key constraint*
 - Specified when the schema is defined.
- ❖ DBMS enforces ICs.
 - Stored data is faithful to real-world meaning.
 - Avoids data entry errors, too!

Primary Key Constraints

- ❖ Key of a relation: **minimal** set of attributes that **uniquely** identify each entity.
 1. No two tuples can have same values in all key fields.
 2. This is not true for any subset of the key.
 - Part 2 false? A *superkey*.
 - If more than 1 key for a relation, *candidate keys*.
 - One of candidate keys is chosen to be the *primary key*.
- ❖ E.g., Students(sid, name, login, age, gpa)

Primary and Candidate Keys in SQL

- ❖ Specify *candidate keys* using **UNIQUE**.
- ❖ Choose one candidate key as the *primary key*.

“For a given student and course,
there is a single grade.”

```
CREATE TABLE Enrolled  
(sid CHAR(20),  
cid CHAR(20),  
grade CHAR(2),  
PRIMARY KEY (sid,cid));
```

“... and no two students in a
course receive the same grade.”

```
CREATE TABLE Enrolled  
(sid CHAR(20),  
cid CHAR(20),  
grade CHAR(2),  
PRIMARY KEY (sid,cid),  
UNIQUE (cid, grade) );
```

Foreign Keys

- ❖ Foreign key: set of fields used to `refer' to the primary key of another relation.
 - Like a `logical pointer' .
- ❖ E.g., Enrolled(*sid*: string, *cid*: string, *grade*: string):
 - *sid* is a foreign key referring to **Students**.

Foreign Keys in SQL

- ❖ Only students listed in the Students relation should be allowed to enroll for courses.

```
CREATE TABLE Enrolled
(sid CHAR(20), cid CHAR(20), grade CHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid) REFERENCES Students );
```

Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

Referential Integrity

- ❖ Referential integrity: any foreign key value must have a matching primary key value in referenced reln.
 - E.g., every *sid* value in Enrolled must appear in **Students**.
 - No dangling references.
- ❖ Can you name a data model w/o referential integrity?

Enforcing Referential Integrity

- ❖ What if an Enrolled tuple with a non-existent student id is inserted?
 - *Reject it!*
- ❖ What if a Students tuple is deleted?
 - **CASCADE**: delete all Enrolled tuples that refer to it.
 - **NO ACTION**: disallow if the Students tuple is referred to.
 - **SET DEFAULT**: set the foreign key to a *default sid*.
 - **SET NULL**: set the foreign key to a special value *null*, denoting '*unknown*' or '*inapplicable*'.
- ❖ Updates to *sid* in Students are treated similarly.

Referential Integrity in SQL

```
CREATE TABLE Enrolled
(sid CHAR(20),
 cid CHAR(20),
 grade CHAR(2),
 PRIMARY KEY (sid,cid),
 FOREIGN KEY (sid)
 REFERENCES Students
 ON DELETE CASCADE
 ON UPDATE NO ACTION);
```


Where do IC's Come From?

- ❖ Based upon real-world business logic.
- ❖ Can check violation in a database instance, but can **NEVER** infer an IC by looking at an instance.
 - An IC is a statement about *all possible* instances!
 - E.g., *name* of the Students relation.

Questions

