



Relational Algebra and Calculus

Yanlei Diao & Michaël Thomazo



Relational Query Languages

- ❖ Relational model allows simple, powerful *querying* of data.
- ❖ Relational query languages:
 - **Declarative**: say “what you want”, not “how to get it”
 - **Formal mathematical model**
 - **Query optimization**

Formal Relational Query Languages

- ❖ Two mathematical languages form the basis for the “real” one, SQL, and for implementation:
 - Relational Algebra: operational, useful for representing execution plans.
 - Relational Calculus: declarative, useful for defining query semantics.



Outline

- ❖ Relational Model
- ❖ Formal Query Languages
 - Relational Algebra
 - Relational Calculus
 - Language Theory

What is an “Algebra”?

- ❖ A mathematical system consisting of:
 - **Operands**: variables or values from which new values can be constructed
 - **Operators**: symbols denoting procedures that construct new values from given values

What is “Relational Algebra”

- ❖ **Relational algebra:**
 - Operands are relations.
 - Operators each take 1 or 2 relations and produce a relation.
- ❖ **Closure property:** relational algebra is closed under the relational model.
 - Relational operators can be arbitrarily *composed*!

Relational Algebra

❖ Basic operations:

- Selection (σ) Selects a subset of rows from a relation.
- Projection (π) Deletes unwanted columns from a relation.
- Cross-product (\times) Allows us to combine two relations.
- Set-difference ($-$) Tuples in reln. 1, but not in reln. 2.
- Union (\cup) Tuples in reln. 1 and in reln. 2.

❖ Additional operations:

- Join (\bowtie), Intersection (\cap), Division ($/$), Renaming (ρ)
- Can be derived from basic operators. Not essential, but useful!

Example Instances

Sailors

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

Reserves

R1

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

Projection

- ❖ Retain only attributes in the *projection list*; delete others.
- ❖ *Schema of result* contains exactly the fields in projection list.

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

sname	rating
yuppy	9
lubber	8
guppy	5
rusty	10

$\pi_{sname, rating}(S2)$

Projection (contd.)

- ❖ Projection operator has to eliminate *duplicates*!
 - Real systems typically don't do duplicate elimination unless the user explicitly asks for it. (Why not?)

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

age
35.0
55.5

$\pi_{age}(S2)$

Selection

- ❖ Select rows that satisfy the *selection condition*; discard others.
- ❖ *Schema of result* identical to schema of input.

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

sid	sname	rating	age
28	yuppy	9	35.0
58	rusty	10	35.0

$$\sigma_{rating > 8}(S2)$$

Selection (contd.)

- ❖ *Composition*: result relation of an operator can be the *input* to another operator.

sid	sname	rating	age
28	yuppy	9	35.0
58	rusty	10	35.0

$$\sigma_{rating > 8}(S2)$$

sname	rating
yuppy	9
rusty	10

$$\pi_{sname, rating}(\sigma_{rating > 8}(S2))$$

Union, Intersection, Set-Difference

- ❖ Set operations:
 - Union (\cup)
 - Intersection (\cap)
 - Set difference ($-$)
- ❖ Take two input relations, which must be *union-compatible*:
 - Same number of fields.
 - Corresponding fields have the same type.
- ❖ What is the *schema* of result?

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

Example Set Operations

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0
44	guppy	5	35.0
28	yuppy	9	35.0

$S1 \cup S2$

Duplicate elimination: remove tuples that have same values in all attributes.

Example Set Operations

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

sid	sname	rating	age
22	dustin	7	45.0

$S1 - S2$

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

sid	sname	rating	age
31	lubber	8	55.5
58	rusty	10	35.0

$S1 \cap S2$

Duplicates?

Cross (Cartesian) Product

❖ $S1 \times R1$: Each row of $S1$ is paired with each row of $R1$.

$S1$	<u>sid</u>	sname	rating	age
	22	dustin	7	45.0
	31	lubber	8	55.5
	58	rusty	10	35.0

$R1$	<u>sid</u>	<u>bid</u>	<u>day</u>
	22	101	10/10/96
	58	103	11/12/96

$S1 \times R1$	(sid)	sname	rating	age	(sid)	bid	day
	22	dustin	7	45.0	22	101	10/10/96
	22	dustin	7	45.0	58	103	11/12/96
	31	lubber	8	55.5	22	101	10/10/96
	31	lubber	8	55.5	58	103	11/12/96
	58	rusty	10	35.0	22	101	10/10/96
	58	rusty	10	35.0	58	103	11/12/96

Cross-Product (contd.)

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

- ❖ *Result schema* inherits all fields of S1 and R1.
 - *Conflict*: Both S1 and R1 have a field called *sid*.
- ❖ Renaming operator: $\rho (C(1 \rightarrow sid1, 5 \rightarrow sid2), S1 \times R1)$

Joins

❖ Condition (theta) Join: $R \bowtie_c S = \sigma_c (R \times S)$

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	58	103	11/12/96

$$S1 \bowtie_{S1.sid < R1.sid} R1$$

- *Result schema* same as that of cross-product.
- But often fewer tuples, more efficient for computation.

Joins

- ❖ Equi-Join: A special case of condition join where the condition θ contains only *equalities*.

sid	sname	rating	age	bid	day
22	dustin	7	45.0	101	10/10/96
58	rusty	10	35.0	103	11/12/96

$$S1 \bowtie_{sid} R1$$

- *Result schema* contains only *one copy* of fields for which equality is specified.
- ❖ Natural Join ($R \bowtie S$): equijoin on *all* common fields.

Example Schema

- ❖ **Sailors**(*sid*: integer, *sname*: string, *rating*: integer, *age*: integer)
- ❖ **Boats**(*bid*: integer, *color*: string)
- ❖ **Reserves**(*sid*: integer, *bid*: integer, *day*: date)

•
•
•
•
•
•
•
Find names of sailors who've reserved boat #103

- ❖ **Sailors**(*sid*: integer, *sname*: string, *rating*: integer, *age*: integer)
- ❖ **Boats**(*bid*: integer, *color*: string)
- ❖ **Reserves**(*sid*: integer, *bid*: integer, *day*: date)

Find names of sailors who've reserved boat #103

❖ Solution 1: $\pi_{sname}((\sigma_{bid=103} Reserves) \bowtie Sailors)$

❖ Solution 2: $\rho(Temp1, \sigma_{bid=103} Reserves)$

$\rho(Temp2, Temp1 \bowtie Sailors)$

$\pi_{sname}(Temp2)$

❖ Solution 3: $\pi_{sname}(\sigma_{bid=103}(Reserves \bowtie Sailors))$

Algebraic equivalence!

Find names of sailors who've reserved a red boat

- ❖ **Sailors**(*sid*: integer, *sname*: string, *rating*: integer, *age*: integer)
- ❖ **Boats**(*bid*: integer, *color*: string)
- ❖ **Reserves**(*sid*: integer, *bid*: integer, *day*: date)

- ❖ Boat color is only available in Boats; so need an extra join:

$$\pi_{sname}((\sigma_{color='red'} Boats) \bowtie Reserves \bowtie Sailors)$$

Find sailors who've reserved a red and a green boat

- ❖ Will a single selection work?
- ❖ Instead, *intersect* sailors who've reserved red boats and sailors who've reserved green boats.

sid is a *key* for Sailors

$\rho (Tempred, \pi_{sid}((\sigma_{color='red'} Boats) \bowtie Reserves))$

$\rho (Tempgreen, \pi_{sid}((\sigma_{color='green'} Boats) \bowtie Reserves))$

$\pi_{sname}((Tempred \cap Tempgreen) \bowtie Sailors)$



Outline

- ❖ Relational Model
- ❖ Formal Query Languages
 - Relational Algebra
 - **Relational Calculus**
 - Language Theory

Relational Calculus

- ❖ Query has the form:

$$\{ \langle x_1, x_2, \dots, x_n \rangle \mid p(\langle x_1, x_2, \dots, x_n \rangle) \}$$

domain variables, or constants

formula

- ❖ Answer includes all *tuples* $\langle x_1, x_2, \dots, x_n \rangle$ that make the formula $p(\langle x_1, x_2, \dots, x_n \rangle)$ *true*.

Formulas

- ❖ *Formula* is recursively defined:
 - *Atomic formulas*: getting tuples from relations, or making comparisons of values
 - *Logical connectives*: \neg , \wedge , \vee
 - *Quantifiers*: \exists , \forall

Free and Bound Variables

- ❖ The use of **quantifiers** $\exists X$ and $\forall X$ in a formula is said to bind X .
 - A variable that is **not bound** is free.
- ❖ Let us revisit the definition of a **query**:

$$\left\{ \langle x_1, x_2, \dots, x_n \rangle \mid p(\langle x_1, x_2, \dots, x_n \rangle) \right\}$$

- ❖ There is an important restriction: the variables **x_1, \dots, x_n** that appear to the left of \mid must be the **only** free variables in the formula $p(\dots)$.

Find sailors rated > 7 who have reserved boat #103

Relational Algebra:

$$\pi_{sname}((\sigma_{bid=103} Reserves) \bowtie (\sigma_{rating>7} Sailors))$$

Relational Calculus:

$$\{X_{sname} \mid \exists \underline{X_{sid}}, X_{rating}, X_{age} \text{ Sailors}(\underline{X_{sid}}, X_{sname}, X_{rating}, X_{age}) \wedge X_{rating} > 7 \\ \wedge \exists X_{bid}, X_{day} \text{ Reserves}(\underline{X_{sid}}, X_{bid}, X_{day}) \wedge X_{bid}=103 \}$$

❖ Where is the join?

- Use \exists to find a tuple in Reserves that 'joins with' the Sailors tuple under consideration.

Find names of sailors who've reserved *all* boats


$$\{X_{\text{sname}} \mid \exists X_{\text{sid}}, X_{\text{rating}}, X_{\text{age}} \langle X_{\text{sid}}, X_{\text{sname}}, X_{\text{rating}}, X_{\text{age}} \rangle \in \text{Sailors} \wedge \\ \forall \langle X_{\text{bid}}, X_{\text{color}} \rangle \in \text{Boats} \\ (\exists X_{\text{day}} \langle X_{\text{sid}}, X_{\text{bid}}, X_{\text{day}} \rangle \in \text{Reserves}) \}$$

❖ To find sailors who've reserved *all red* boats:

$$\{X_{\text{sname}} \mid \exists X_{\text{sid}}, X_{\text{rating}}, X_{\text{age}} \langle X_{\text{sid}}, X_{\text{sname}}, X_{\text{rating}}, X_{\text{age}} \rangle \in \text{Sailors} \wedge \\ \forall \langle X_{\text{bid}}, X_{\text{color}} \rangle \in \text{Boats} \\ (X_{\text{color}} \neq \text{'red'} \vee \exists X_{\text{day}} \langle X_{\text{sid}}, X_{\text{bid}}, X_{\text{day}} \rangle \in \text{Reserves}) \}$$

Find names of sailors who've reserved all boats

$$\{X_{\text{sname}} \mid \exists X_{\text{sid}}, X_{\text{rating}}, X_{\text{age}} \langle X_{\text{sid}}, X_{\text{sname}}, X_{\text{rating}}, X_{\text{age}} \rangle \in \text{Sailors} \wedge \\ \forall \langle X_{\text{bid}}, X_{\text{color}} \rangle \in \text{Boats} \\ (\exists X_{\text{day}} \langle X_{\text{sid}}, X_{\text{bid}}, X_{\text{day}} \rangle \in \text{Reserves}) \quad \}$$



$$\{X_{\text{sname}} \mid \exists X_{\text{sid}}, X_{\text{rating}}, X_{\text{age}} \langle X_{\text{sid}}, X_{\text{sname}}, X_{\text{rating}}, X_{\text{age}} \rangle \in \text{Sailors} \wedge \\ \neg \exists \langle X_{\text{bid}}, X_{\text{color}} \rangle \in \text{Boats} \\ (\neg \exists X_{\text{day}} \langle X_{\text{sid}}, X_{\text{bid}}, X_{\text{day}} \rangle \in \text{Reserves}) \quad \}$$

Find the names of sailors who've reserved all boats

- ❖ Step 1: for each sailor, check if there exists a boat that he has not reserved (called formula F).

$$\rho(S_neg, \pi_{sid}((\pi_{sid} Reserves) \times (\pi_{bid} Boats) - (\pi_{sid,bid} Reserves)))$$

- ❖ Step 2: find sailors for which F is not true and retrieve their names

$$\pi_{sname}((\pi_{sid} Reserves - S_neg) \bowtie Sailors)$$

‘-’ : the only way to express negation in relational algebra!



Outline

- ❖ Relational Model
- ❖ Formal Query Languages
 - Relational Algebra
 - Relational Calculus
 - **Language Theory**

Query Languages

- ❖ The three languages we consider:
 - Relational Algebra (RA)
 - Relational Calculus (RC)
 - Structured Query Language (SQL)

Find sailors rated > 7 who have reserved boat #103

Relational Algebra:

$\pi_{sname}((\sigma_{bid=103} Reserves) \bowtie (\sigma_{rating>7} Sailors))$

Relational Calculus:

$\{X_{sname} \mid \exists X_{sid}, X_{rating}, X_{age} Sailors(X_{sid}, X_{sname}, X_{rating}, X_{age}) \wedge X_{rating} > 7$
 $\wedge \exists X_{bid}, X_{day} Reserves(X_{sid}, X_{bid}, X_{day}) \wedge X_{bid} = 103 \}$

SQL:

SELECT sname
FROM Sailors S, Reserves R
WHERE S.sid=R.sid and s.rating>7 and R.bid = '103';

3. Final projection

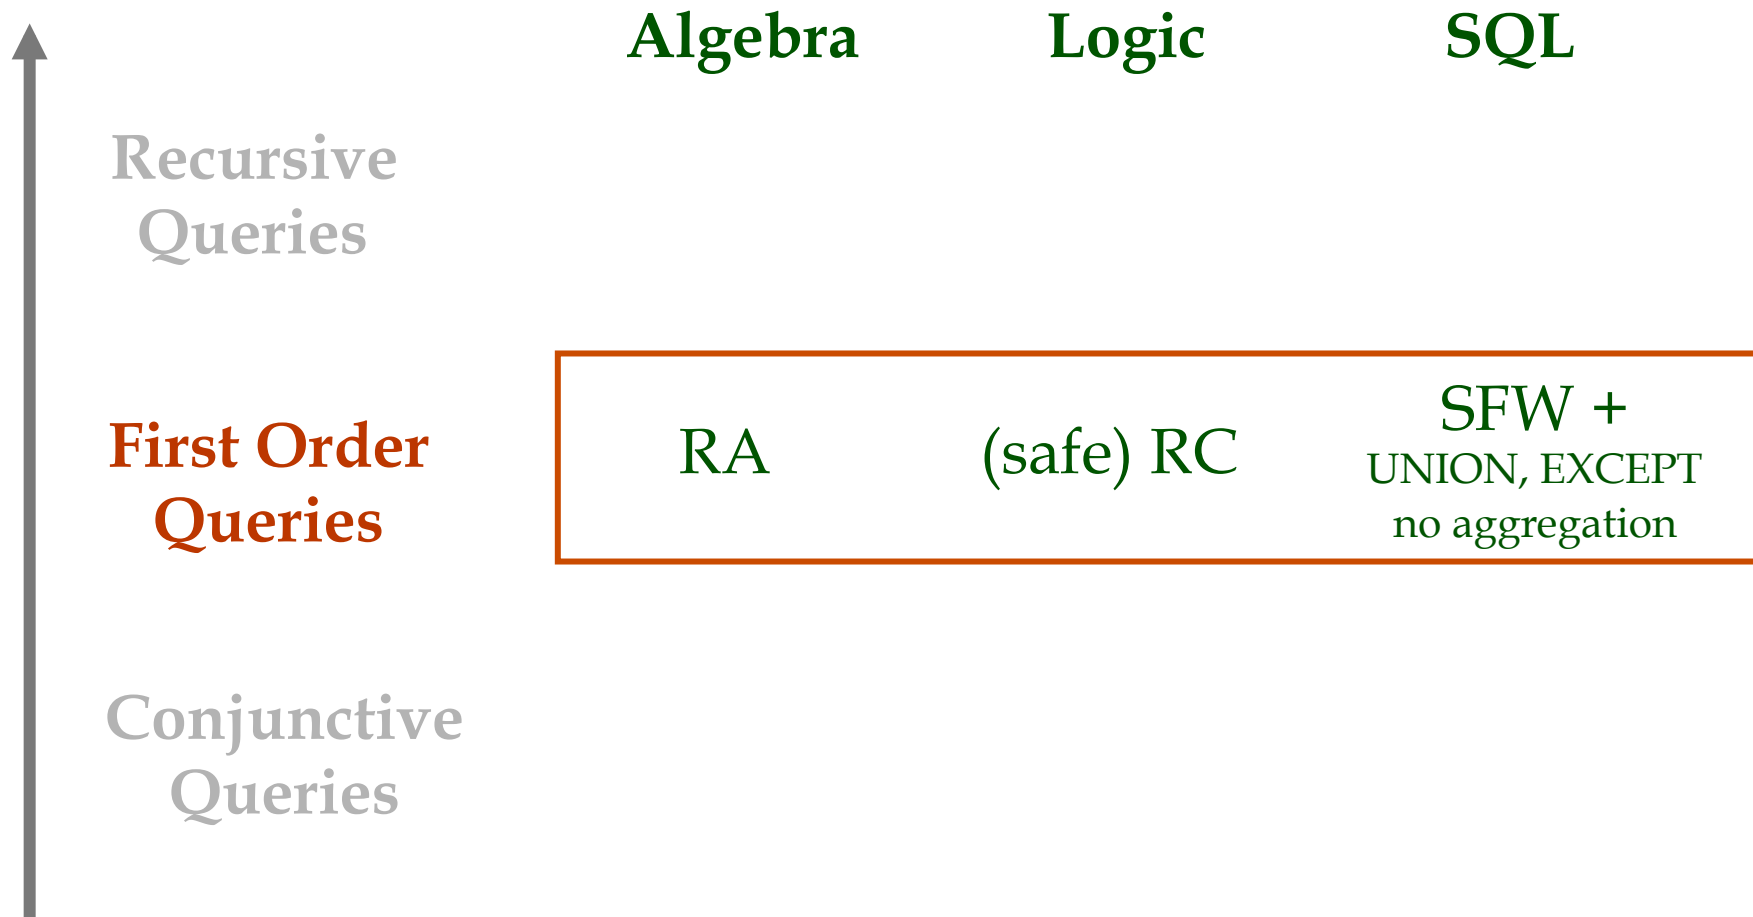
1. Implicit cross product

2. Selection on the results of cross product

Unsafe Queries, Expressive Power

- ❖ *Unsafe* queries in calculus: some queries can have an infinite number of answers.
 - e.g., $\{S \mid \neg (S \in Sailors)\}$
- ❖ Equivalence between RA and Safe RC: every query that can be expressed in *relational algebra* can be expressed as a *safe query in relational calculus*; the converse is also true.
- ❖ SQL can express every query that is expressible in relational algebra/calculus.

Query Language Classes



Query Language Classes


	Algebra	Logic	SQL
Recursive Queries			
First Order Queries	RA	(safe) RC	SFW + UNION, EXCEPT no aggregation
Conjunctive Queries	RA: σ, π, \times	Single datalog rule	S ^d FW no aggregation

Conjunctive Queries (CQ)

- ❖ A **subset** of FO queries (i.e., less expressive).
- ❖ CQs have “better” theoretical properties than arbitrary queries.
- ❖ Query optimizer handles CQs the best--it tries to
 - flatten a nested query to a single CQ
 - break a large query into many CQs

CQ in Rule-based (Datalog) Notation

Variables



- ❖ R : Extensional database (EDB) -- stored
- ❖ P : Intentional database (IDB) -- computed

Find sailors rated > 7 who have reserved boat #103

$$\mathbf{P}(x,z) \leftarrow \mathbf{R}_1(x, y) \ \& \ \mathbf{R}_2(y, z)$$

$$P(X_{\text{sname}}) \leftarrow \text{Sailors}(X_{\text{sid}}, X_{\text{sname}}, X_{\text{rating}}, X_{\text{age}}) \ \& \ X_{\text{rating}} > 7 \\ \& \text{Reserves}(X_{\text{sid}}, 103, X_{\text{day}})$$

Properties of CQs

❖ Satisfiability

- A query q is *satisfiable* if there exists at least one database instance D such that $q(D)$ is non-empty.
- Theorem: *Every CQ is satisfiable.*

❖ Monotonicity

- A query Q is *monotonic* if for two database instances $D1$ and $D2$, $D1 \subseteq D2$ implies $Q(D1) \subseteq Q(D2)$.
- Theorem: *Every CQ is monotonic.*

Beyond First-Order Queries

	Algebra	Logic	SQL
Recursive Queries	?		
First Order Queries	RA	(safe) RC	SFW + UNION, EXCEPT no aggregation
Conjunctive Queries	RA: σ, π, \times	Single datalog rule	S ^d FW no aggregation

Limitation of FO Queries

- ❖ Let $D = \{E(x,y)\}$ represent a graph
- ❖ Query $\text{path}(x,z) =$
 - all x,z such that there is a path from x to z .
- ❖ **Theorem:** $\text{path}(x,z)$ *cannot* be expressed in FO.

Find all of Mary's ancestors

ParentOf

Parent	Child
Mike	Joe
Joe	Alice
Joe	Bob
Alice	Mary
...	...

❖ Can you write a query in SQL?

SQL with Recursion

Relation to be computed recursively

```
WITH RECURSIVE Ancestor(anc, desc) AS
((SELECT parent AS anc, child AS desc
  FROM ParentOf)
 UNION
 ((SELECT A.anc, p.child AS desc
   FROM Ancestor A, ParentOf P
   WHERE A.desc = P.parent)
 )
 )
 SELECT anc
 FROM Ancestor
 WHERE desc = 'Mary' ;
```

← Base case

← Recursion

☐ ☐ ☐ ☐ ☐ ☐

ParentOf

Iter 3: $\{ \}$

Parent	Child
Mike	Joe
Joe	Alice
Joe	Bob
Alice	Mary

Query result:

Alice	Mary
Joe	Mary
Mike	Mary

Recursion in Datalog

$\text{Ancestor}(x,y) \leftarrow \text{ParentOf}(x,y)$

$\text{Ancestor}(x,z) \leftarrow \text{Ancestor}(x,y) \ \& \ \text{ParentOf}(y,z)$

Use of IDB in Body

Implicit UNION

$\text{AncestorOfMary}(x) \leftarrow \text{Ancestor}(x, \text{'Mary'})$

A More Complete Picture

	Algebra	Logic	SQL
Recursive Queries	Fixed point operator	Datalog (recursion)	Full SQL (recursion)
First Order Queries	RA	(safe) RC	SFW + UNION, EXCEPT no aggregation
Conjunctive Queries	RA: σ, π, \times	Single Datalog rule	S ^d FW no aggregation

Questions

