



---

# *Structured Query Language*

---

Yanlei Diao

Slides Courtesy of R. Ramakrishnan, J. Gehrke, and G. Miklau



# *Structured Query Language (SQL)*

- ❖ *Data Manipulation Language (DML)*
  - posing queries, operating on tuples
- ❖ *Data Definition Language (DDL)*
  - operating on tables/views
- ❖ Extension from Relational Algebra / Calculus
  - From a set to a **multi-set** (bag) based model
  - Extending first order expressive power with **aggregation**, and **recursion**

# SQL Overview

---

## ❖ Query capabilities

- SELECT-FROM-WHERE blocks
- Set operations (union, intersect, except)
- Nested queries (correlation)
- Aggregation and grouping
- Ordering
- Null values

## ❖ Database updates

## ❖ Tables and views

## ❖ Integrity constraints

## Example Instances

*S1*

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

*S2*

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

*R1*

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

# Basic SQL Query

SELECT	[ <b>DISTINCT</b> ] <i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i> ;

- ❖ *relation-list*: a list of input relation names, possibly each with a *range-variable*.
- ❖ *qualification*: *predicates* combined using AND, OR and NOT.
  - predicate: attr op const or attr1 op attr2, *op* is <, >, >=, <=, =, <>
- ❖ *target-list*: a list of attributes to display in output
  - **DISTINCT** indicates no duplicates in the answer. Default is that duplicates are not eliminated!

# Conceptual Evaluation Strategy

SELECT	[DISTINCT] <i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i> ;

- ❖ *relation-list*: cross-product (  $\times$  )
- ❖ *qualification*: selection (  $\sigma$  )
- ❖ *target-list*: projection (  $\pi$  )
  - duplicate elimination if DISTINCT
- ❖ This is possibly the least efficient way to execute the query! Leave the issue to Query Optimization...

# Example of Conceptual Evaluation

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid=R.sid AND R.bid=103;
```

(sid)	sname	rating	age	(sid)	bid	day	
22	dustin	7	45.0	22	101	10/10/96	X
22	dustin	7	45.0	58	103	11/12/96	X
31	lubber	8	55.5	22	101	10/10/96	X X
31	lubber	8	55.5	58	103	11/12/96	X
58	rusty	10	35.0	22	101	10/10/96	X X
58	rusty	10	35.0	58	103	11/12/96	

☛ What is the relational algebra for this query?

## *Relational Algebra for the Query*

```
SELECT S.sname  
FROM   Sailors S, Reserves R  
WHERE  S.sid=R.sid AND R.bid=103;
```

$$\pi_{sname}((\sigma_{bid=103} Reserves) \bowtie Sailors)$$



## *A Note on Range Variables*

- ❖ Really needed only if the same relation appears twice in the FROM clause.

```
SELECT sname
FROM   Sailors, Reserves
WHERE  Sailors.sid=Reserves.sid
      AND bid=103;
```

OR

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid=R.sid AND bid=103;
```

*It is good style,  
however, to use  
range variables  
always!*

•  
•  
•  
•  
•  
•  
•  
Find sailors who've reserved at least one boat

```
SELECT S.sid  
FROM   Sailors S, Reserves R  
WHERE  S.sid=R.sid;
```

- Would adding DISTINCT to this query change the answer set?
- What if we replace *S.sid* by *S.sname* in the SELECT clause and then add DISTINCT?

# String Pattern Matching

```
SELECT S.age
FROM   Sailors S
WHERE  S.sname LIKE 'A_%M';
```

- ❖ *Find the ages of sailors whose names begin with 'A', end with 'M', and contain at least three characters.*
- ❖ **LIKE** is used for string matching.
  - ``_`` stands for any one character.
  - ``%`` stands for 0 or more arbitrary characters.

# Arithmetic Expressions

---

```
SELECT S.age, age1 = S.age-5, 2*S.age AS age2
FROM   Sailors S
WHERE  S.sname LIKE 'A%M' ;
```

- ❖ *For sailors whose names begin with 'A' and end with 'M', return triples (of ages of sailors and two fields defined by expressions)*
- ❖ Arithmetic expressions create derived attributes in SELECT.
  - **AS** and **=** are two ways to name fields in the result.
- ❖ They can also appear in the predicates in WHERE.

# SQL Overview

---

- ❖ Query capabilities

- SELECT-FROM-WHERE blocks
- Set operations (union, intersect, except)
- Nested queries (correlation)
- Aggregation and grouping
- Ordering
- Null values

- ❖ Database updates

- ❖ Tables and views

- ❖ Integrity constraints

Find sid's of sailors who've reserved a red or a green boat

---

- ❖ If we replace **OR** by **AND** in this query, what do we get?

```
SELECT DISTINCT R.sid
FROM Reserves R, Boats B
WHERE R.bid=B.bid
      AND (B.color= 'red' OR B.color= 'green' )
```

- ❖ **UNION**: computes the union of any two *union-compatible* sets of tuples (which are themselves the result of SQL queries).

```
SELECT DISTINCT S.sid
FROM Reserves R, Boats B
WHERE R.bid=B.bid
      AND B.color='red'
UNION
SELECT DISTINCT S.sid
FROM Reserves R, Boats B
WHERE R.bid=B.bid
      AND B.color= 'green';
```

Find sid's of sailors who've reserved a red and a green boat

---

- ❖ **INTERSECT**: computes the intersection of any two *union-compatible* sets of tuples.

```
SELECT DISTINCT S.sid
FROM Reserves R, Boats B
WHERE R.bid=B.bid
      AND B.color= 'red'
INTERSECT
SELECT DISTINCT S.sid
FROM Reserves R, Boats B
WHERE R.bid=B.bid
      AND B.color= 'green' ;
```

Need **DISTINCT** to be equivalent!

```
SELECT DISTINCT S.sid
FROM Reserves R1, Boats B1,
      Reserves R2, Boats B2
WHERE R1.bid=B1.bid AND R2.bid=B2.bid
      AND (B1.color= 'red' AND B2.color= 'green' )
      AND R1.sid=R2.sid;
```

•  
•  
•  
•  
•  
•  
•  
•  
•  
•  
*Find sid's of sailors who've reserved ...*

---

- ❖ Also available: **EXCEPT** (What does this query return?)

```
SELECT DISTINCT S.sid  
FROM Reserves R, Boats B  
WHERE R.bid=B.bid  
      AND B.color= 'red'
```

**EXCEPT**

```
SELECT DISTINCT S.sid  
FROM Reserves R, Boats B  
WHERE R.bid=B.bid  
      AND B.color= 'green' ;
```



# SQL Overview

---

- ❖ Query capabilities
  - SELECT-FROM-WHERE blocks
  - Set operations (union, intersect, except)
  - Nested queries (correlation)
  - Aggregation & Grouping
  - Ordering
  - Null values
- ❖ Database updates
- ❖ Tables and views
- ❖ Integrity constraints

# Nested Queries

---

- ❖ A **nested query** has another query embedded within it.
- ❖ The embedded query is called the **subquery**.
- ❖ The subquery often appears in the WHERE clause:

```
SELECT  S.sname
FROM    Sailors S
WHERE   S.sid IN ( SELECT R.sid
                  FROM Reserves R
                  WHERE R.bid = 103 );
```

- ❖ Subqueries are also possible in the FROM clause.

# Conceptual Evaluation, extended

- ❖ For each row in the cross-product of the outer query, evaluate the WHERE condition by *re-computing the subquery*.

```
SELECT  S.sname
FROM    Sailors S
WHERE   S.sid IN ( SELECT R.sid
                  FROM Reserves R
                  WHERE R.bid = 103 );
```

However, this query is equivalent to (can be *simplified* to):

```
SELECT  S.sname
FROM    Sailors S, Reserves R
WHERE   S.sid=R.sid AND R.bid=103;
```

# Correlated Subquery

- ❖ A subquery that depends on the table(s) mentioned in the outer query is a **correlated subquery**.
- ❖ In conceptual evaluation, must recompute subquery for each row of the outer query.

```
SELECT S.sname
FROM Sailors S
WHERE EXISTS ( SELECT *
                FROM Reserves R
                WHERE R.bid = 103
                  AND R.sid = S.sid );
```

Correlation



# Set Comparison Operators in WHERE

- ❖ Set comparison, optionally with a proceeding **NOT**:
  - **EXISTS**  $R$  -- true if  $R$  is non-empty
  - $attr$  **IN**  $R$  -- true if  $R$  contains  $attr$
  - **UNIQUE**  $R$  -- true if no duplicates in  $R$
- ❖ Arithmetic operator  $op \{<, <=, =, < >, >=, >\}$  and **ALL/ANY**:
  - $attr op$  **ALL**  $R$  -- all elements of  $R$  satisfy condition
  - $attr op$  **ANY**  $R$  -- some element of  $R$  satisfies condition

$'attr$  **IN**  $R$ '                      equivalent to                       $'attr =$  **ANY**  
 $R$ '

$'attr$  **NOT IN**  $R$ '                      equivalent to                       $'attr < >$  **ALL**  
 $R$ '

# *Finding Extreme Values*

---

- ❖ Find the sailors with the *highest* rating

```
SELECT S.sid  
FROM   Sailors S  
WHERE  S.rating >= ALL ( SELECT S2.rating  
                        FROM Sailors S2 );
```

# Please Write SQL

---

- ❖ Find sailors whose rating is higher than *some* sailor named Harry.

```
SELECT S.sid
FROM   Sailors S
WHERE  S.rating > ANY (SELECT S2.rating
                      FROM Sailors S2
                      WHERE S2.name = 'Harry');
```

- ❖ Find sailors whose rating is higher than *all* sailors named Harry.

```
SELECT S.sid
FROM   Sailors S
WHERE  S.rating > ALL (SELECT S2.rating
                      FROM Sailors S2
                      WHERE S2.name = 'Harry');
```

Find sailors who've reserved *all* boats.

---

(1) SELECT S.sname  
FROM Sailors S  
WHERE NOT EXISTS (

SELECT B.bid  
FROM Boats B  
WHERE NOT EXISTS (  
SELECT R.bid  
FROM Reserves R  
WHERE R.bid=B.bid  
AND R.sid=S.sid));

(2)

SELECT S.sname  
FROM Sailors S  
WHERE NOT EXISTS  
((SELECT B.bid  
FROM Boats B)  
EXCEPT  
(SELECT R.bid  
FROM Reserves R  
WHERE R.sid=S.sid));



# SQL Overview

---

- ❖ Query capabilities
  - SELECT-FROM-WHERE blocks
  - Set operations (union, intersect, except)
  - Nested queries (correlation)
  - Aggregation and grouping
  - Ordering
  - Null values
- ❖ Database updates
- ❖ Tables and views
- ❖ Integrity constraints

# Example Aggregate Operators

```
SELECT COUNT(*)  
FROM Sailors S;
```

```
SELECT AVG(S.age)  
FROM Sailors S  
WHERE S.rating=10;
```

```
SELECT AVG(DISTINCT S.age)  
FROM Sailors S  
WHERE S.rating=10;
```

```
SELECT COUNT(DISTINCT S.rating)  
FROM Sailors S  
WHERE S.sname= 'Bob' ;
```

```
SELECT S.sname  
FROM Sailors S  
WHERE S.rating=  
      (SELECT MAX(S2.rating)  
       FROM Sailors S2);
```

# Aggregate Operators

COUNT (\*)

COUNT ([DISTINCT] A)

SUM ([DISTINCT] A)

AVG ([DISTINCT] A)

MAX (A)

MIN (A)

*multiple columns*

*single column*

- ❖ Take a relation (single column or multiple columns), return a **value**.
- ❖ Significant extension of relational algebra.

## *Find name and age of the oldest sailor(s)*

---

- ❖ The first query is illegal! (We'll look into the reason a bit later, when we discuss **GROUP BY**.)

```
SELECT S.sname, MAX (S.age)
FROM Sailors S;
```

```
SELECT S.sname, S.age
FROM Sailors S
WHERE S.age =
      (SELECT MAX (S2.age)
       FROM Sailors S2);
```

# Motivation for Grouping

---

- ❖ What if we want to apply aggregate operators to each *group* (subset) of tuples?
- ❖ Find the age of the youngest sailor for *each* rating level.
  - If we know that rating values  $\in [1, 10]$ , write 10 queries like:

For  $i = 1, 2, \dots, 10$ :

```
SELECT MIN (S.age)
FROM Sailors S
WHERE S.rating = i
```

- In general, we don't know how many rating levels exist, and what the rating values for these levels are!

# Queries with GROUP BY and HAVING

```
SELECT      [DISTINCT] target-list
FROM        relation-list
WHERE       qualification
GROUP BY    grouping-list
[HAVING     group-qualification];
```

- ❖ A *group* is a set of tuples that have the same value for all attributes in *grouping-list*.
- ❖ Query returns a *single* answer tuple for each group!
- ❖ The *target-list* can only contain:
  - (i) attributes in the *grouping-list* (e.g., *S.rating*), or
  - (ii) *aggregate operations* on other attributes, e.g., MIN (*S.age*).

# Conceptual Evaluation, extended

---

- ❖ The cross-product of *relation-list* is computed.
- ❖ Tuples that fail *qualification* are discarded.
- ❖ The remaining tuples are partitioned into groups by the value of attributes in *grouping-list*.
- ❖ The *group-qualification*, if present, eliminates some groups.
  - *Group-qualification* must have a **single** value **per group**!
- ❖ A **single** answer tuple is produced for each qualifying group.

*Find age of the youngest sailor with age  $\geq 18$ ,  
for each rating with at least 2 such sailors*

```
SELECT S.rating, MIN (S.age)
           AS minage
FROM Sailors S
WHERE S.age  $\geq$  18
GROUP BY S.rating
HAVING COUNT (*) > 1;
```

*Answer relation:*

rating	minage
3	25.5
7	35.0
8	25.5

*Sailors instance:*

sid	sname	rating	age
22	dustin	7	45.0
29	brutus	1	33.0
31	lubber	8	55.5
32	andy	8	25.5
58	rusty	10	35.0
64	horatio	7	35.0
71	zorba	10	16.0
74	horatio	9	35.0
85	art	3	25.5
95	bob	3	63.5
96	frodo	3	25.5




*Find age of the youngest sailor with age  $\geq 18$ ,  
for each rating with at least 2 such sailors.*

rating	age
7	45.0
1	33.0
8	55.5
8	25.5
10	35.0
7	35.0
10	16.0
9	35.0
3	25.5
3	63.5
3	25.5



rating	age
1	33.0
3	25.5
3	63.5
3	25.5
7	45.0
7	35.0
8	55.5
8	25.5
9	35.0
10	35.0



rating	minage
3	25.5
7	35.0
8	25.5

•  
•  
•  
•  
•  
•  
•  
*Find those ratings for which the average age is  
the minimum over all ratings*

---

```
SELECT Temp.rating, Temp.avgage
FROM (SELECT S.rating, AVG (S.age) AS avgage
      FROM Sailors S
      GROUP BY S.rating) AS Temp
WHERE Temp.avgage = (SELECT MIN (Temp.avgage)
                    FROM Temp);
```

- ❖ **Derived table:** result of an SQL query as input to the FROM clause of another query
  - Computed once before the other query is evaluated.

# SQL Overview

---

## ❖ Query capabilities

- SELECT-FROM-WHERE blocks
- Set operations (union, intersect, except)
- Nested queries (correlation)
- Aggregation & Grouping
- Ordering
- Null values

## ❖ Database updates

## ❖ Tables and views

## ❖ Integrity constraints

# ORDER BY

---

- ❖ Return the name and age of sailors rated level 8 or above *in increasing (decreasing) order of age*.

```
SELECT S.sname, S.age  
FROM Sailors S  
WHERE S.rating > 8  
ORDER BY S.age [ASC | DESC];
```

## TOP-K Queries

---

- ❖ Return the name and age of the *ten youngest* sailors rated level 8 or above.

```
SELECT S.sname, S.age  
FROM Sailors S  
WHERE S.rating >= 8  
ORDER BY S.age ASC  
LIMIT 10;
```

# SQL Overview

---

## ❖ Query capabilities

- SELECT-FROM-WHERE blocks
- Set operations (union, intersect, except)
- Nested queries (correlation)
- Aggregation & Grouping
- Ordering
- Null values

## ❖ Database updates

## ❖ Tables and views

## ❖ Integrity constraints

# *NULL Values in SQL*

---

- ❖ Whenever we don't have a value, put a **NULL**.
- ❖ Can mean many things:
  - Value does not exist
  - Value exists but is unknown
  - Value not applicable
- ❖ The schema specifies for each attribute whether it can be null (e.g., **NOT NULL**)
- ❖ How does SQL cope with tables that have NULLs ?

# *Null Values*

---

- ❖ If  $x = \text{NULL}$ , then  $4 \cdot (3 - x) / 7$  is still **NULL**
- ❖ If  $x = \text{NULL}$ , then  $x = \text{"Joe"}$  is **UNKNOWN**
- ❖ In SQL there are three boolean values:
  - $\text{FALSE} = 0$
  - $\text{UNKNOWN} = 0.5$
  - $\text{TRUE} = 1$



# *Coping with Unknown Values*

- ❖  $C1 \text{ AND } C2 = \min(C1, C2)$
- ❖  $C1 \text{ OR } C2 = \max(C1, C2)$
- ❖  $\text{NOT } C1 = 1 - C1$

```
SELECT *  
FROM   Person  
WHERE  (age < 25) AND  
        (height > 6 OR weight > 190);
```

E.g.  
age=20  
height=NULL  
weight=200

- ❖ Rule in SQL: include only tuples that yield **TRUE**

# *Anomaly Associated with Null's*

- ❖ Unexpected behavior:

```
SELECT *  
FROM   Person  
WHERE  age < 25 OR age >= 25;
```

E.g.  
John's age is **NULL**

- ❖ Some person is not included!

# Null Values

---

- ❖ Can test for NULL explicitly:
  - x IS NULL
  - x IS NOT NULL

```
SELECT *  
FROM   Person  
WHERE  age < 25 OR age >= 25 OR age IS NULL;
```

- ❖ Now it includes all people.

# SQL Overview

---

- ❖ Query capabilities

- SELECT-FROM-WHERE blocks
- Set operations (union, intersect, except)
- Nested queries (correlation)
- Aggregation & Grouping
- Ordering
- Null values

- ❖ Database updates

- ❖ Tables and views

- ❖ Integrity constraints

# *Modifying the Database*

---

Three kinds of modifications:

- ❖ Insert - create new tuple(s)
- ❖ Delete - remove existing tuple(s)
- ❖ Update - modify existing tuple(s)
  
- ❖ Sometimes they are all called “updates”.

# *Insertions*

---

General form:

```
INSERT INTO R(A1,..., An)
VALUES (v1,..., vn);
```

Example: Insert a new sailor to the database:

```
INSERT INTO Sailors(sid, sname, rating, age)
VALUES (3212, 'Fred' , 9, 44);
```

Can omit attributes; a missing attribute is NULL.  
May drop attribute names if give values of all attributes in order.

# Insertions

---

Example: Insert *multiple* tuples to Sailors:

```
INSERT INTO Sailors(sid, sname)

SELECT B.id, B.name
FROM   Boaters B
WHERE  Boaters.rank = 'captain';
```

The query replaces the VALUES keyword.

# Deletions

---

Example: delete all tuples that satisfy a condition

```
DELETE  
FROM Sailors  
WHERE S.sname = 'Harry';
```

Fact about SQL: there is no way to delete only a single occurrence of a tuple that appears twice in a relation.



# Updates

Examples:

```
UPDATE Employees  
SET salary = salary * 1.1;
```

```
UPDATE Sailors S  
SET S.rating = s.rating + 1  
WHERE S.sid IN  
      (SELECT sid  
       FROM Reserves R  
       WHERE R.date = 'Oct, 25');
```

# SQL Overview

---

- ❖ Query capabilities
  - SELECT-FROM-WHERE blocks
  - Set operations (union, intersect, except)
  - Nested queries (correlation)
  - Aggregation & Grouping
  - Ordering
  - Null values
- ❖ Database updates
- ❖ Tables and views
- ❖ Integrity constraints

# Creating Tables

```
CREATE TABLE Sailors
(  sid INTEGER,
   sname CHAR(50) NOT NULL,
   rating INTEGER,
   age REAL,
   PRIMARY KEY (sid));
```

```
CREATE TABLE Boats
(  bid INTEGER,
   bname CHAR (20),
   color CHAR(20),
   PRIMARY KEY (bid)
   UNIQUE (bname));
```

```
CREATE TABLE Reserves
(  sid INTEGER,
   bid INTEGER,
   day DATE,
   PRIMARY KEY (sid,bid,day),
   FOREIGN KEY (sid) REFERENCES Sailors
       ON DELETE NO ACTION ON UPDATE CASCADE
   FOREIGN KEY (bid) REFERENCES Boats
       ON DELETE SET DEFAULT ON UPDATE CASCADE);
```

# *Destroying and Altering Tables*

```
DROP TABLE Sailors;
```

- ❖ Destroys the Sailors relation, including schema and data.

```
ALTER TABLE Sailors  
    ADD COLUMN credit_card:CHAR(40);
```

- ❖ The schema is altered by adding a new field; every tuple in the current instance is extended with a *null* value in the new field.

# Views

- ❖ A view is like a relation, but we store a *definition*, rather than a set of tuples.

```
CREATE VIEW RedBoatLovers (sid, name, bid)
  AS SELECT  S.sid, S.sname, B.bid
    FROM    Sailors S, Reserves R, Boats B
    WHERE   S.sid = R.sid and R.bid = B.bid
           and B.color= 'red' ;
```

- ❖ Views can be dropped using **DROP VIEW** command.
  - **DROP TABLE** if there's a view on the table?

# Uses of Views

---

- ❖ Views can be used to present necessary information (or a summary), while *hiding details in underlying relation(s)*.
- ❖ Security/Privacy
  - E.g., hiding sailors' credit card from the boat repair dept.
- ❖ Logical data independence
  - User application defined on a view is unchanged when underlying table changes
- ❖ Computational benefits
  - Result of a complex query is frequently used; materialize it.
  - Online Analytical Processing (OLAP)

# *SQL Overview*

---

- ❖ Query capabilities
  - SELECT-FROM-WHERE blocks
  - Set operations (union, intersect, except)
  - Nested queries (correlation)
  - Aggregation and grouping
  - Ordering
  - Null values
- ❖ Database updates
- ❖ Tables and views
- ❖ Integrity constraints

# *Integrity Constraints (Review)*

---

- ❖ Types of integrity constraints in SQL:
  - *Attribute constraints*: domain, NOT NULL
  - *Key constraints*: PRIMARY KEY, UNIQUE
  - *Foreign key constraints*: FOREIGN KEY
  - *General constraints*: CHECK, ASSERTION
- ❖ Inserts/deletes/updates that violate IC's are disallowed.



# General Constraints

- ❖ Two forms: **CHECK** (single table constraint) and **ASSERTION** (multiple-table constraint).

```
CREATE TABLE Sailors
(  sid INTEGER,
   sname CHAR(50),
   rating INTEGER,
   age REAL,
   PRIMARY KEY (sid),
   CHECK ( rating >= 1
           AND rating <= 10));
```

# Constraints over Multiple Relations

*Number of boats plus number of sailors is < 100:*

```
CREATE ASSERTION smallClub  
CHECK  
( (SELECT COUNT (S.sid) FROM Sailors S) +  
  (SELECT COUNT (B.bid) FROM Boats B) < 100 );
```

- ❖ ASSERTION is a constraint over both tables; checked whenever one of the table is modified.

# Questions

---

