



---

# *Relational Query Optimization*

---

Yanlei Diao



# Overview of Query Evaluation

---

- ❖ *Query evaluation plan*: tree of *relational algebra* operators, with choice of algorithm for each operator.
- ❖ *Query optimization*: given a query, many plans are possible
  - Ideally, find the most efficient plan.
  - In practice, avoid worst plans in practice.

# *Outline of topics*


---

- ❖ Query plans and equivalences
- ❖ Query optimization issues
  - Plan space
  - Cost estimation
  - Plan search
- ❖ Handling nested queries

# Relational Algebra Tree

```
SELECT S.sname
FROM   Reserves R, Sailors S
WHERE  R.sid=S.sid AND
       R.bid=100 AND S.rating>5
```

*Relational Algebra Tree:*



*Expression in Relational Algebra (RA):*

$\pi_{\text{sname}} (\sigma_{\text{bid}=100 \wedge \text{rating} > 5} (\text{Reserves} \bowtie_{\text{sid}=\text{sid}} \text{Sailors}))$

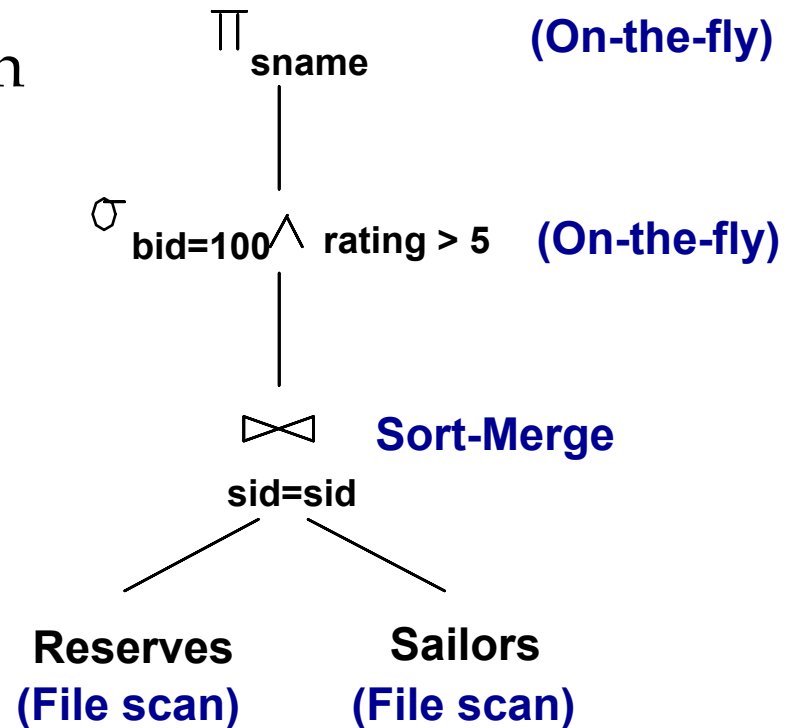
# Query Evaluation Plan

❖ *Query evaluation plan* extends an RA tree with:

- access method for each relation;
- implementation method for each other operator.

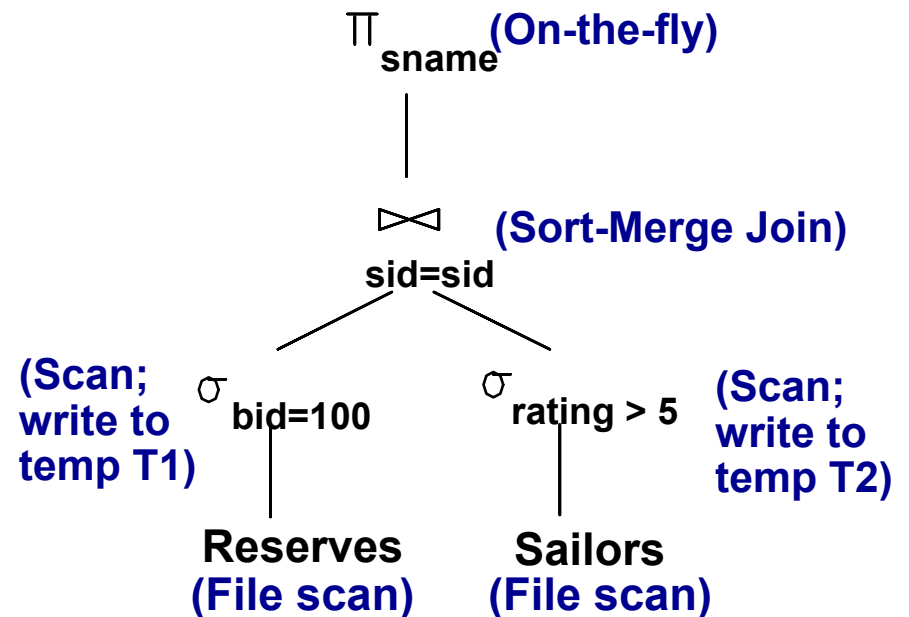
❖ What are the missed opportunities?

- Selections could have been 'pushed' earlier.
- Use of indexes.
- More efficient joins.



# Query Plan 1 (Selection Pushed Down)

- ❖ *Push selections below the join.*
- ❖ *Materialization vs. Pipelining:*
  - Store a temporary relation T, if the subsequent join needs to scan T multiple times.
  - The opposite is pipelining.



- ❖ With 5 buffer pages, cost of plan:
  - Scan Reserves (1000) + write temp T1 (10 pages, if we have 100 boats, uniform distribution).
  - Scan Sailors (500) + write temp T2 (250 pages, if we have 10 ratings).
  - **Sort-Merge join**: Sort T1 ( $2 \times 2 \times 10$ ), sort T2 ( $2 \times 4 \times 250$ ), merge (10+250).
  - Total = 4060 page I/Os.

# Indexes

---

- ❖ A **tree** index *matches* (a conjunction of) terms if the attributes in the terms form a *prefix* of the search key.
  - Tree index on  $\langle a, b, c \rangle$
  - $a=5$  AND  $b=3$  ?
  - $a=5$  AND  $b>6$  ?
  - $b=3$  ?

# Query Plan 3 (Using Indexes)

- ❖ *Selection using index*: clustered index on *bid* of Reserves.

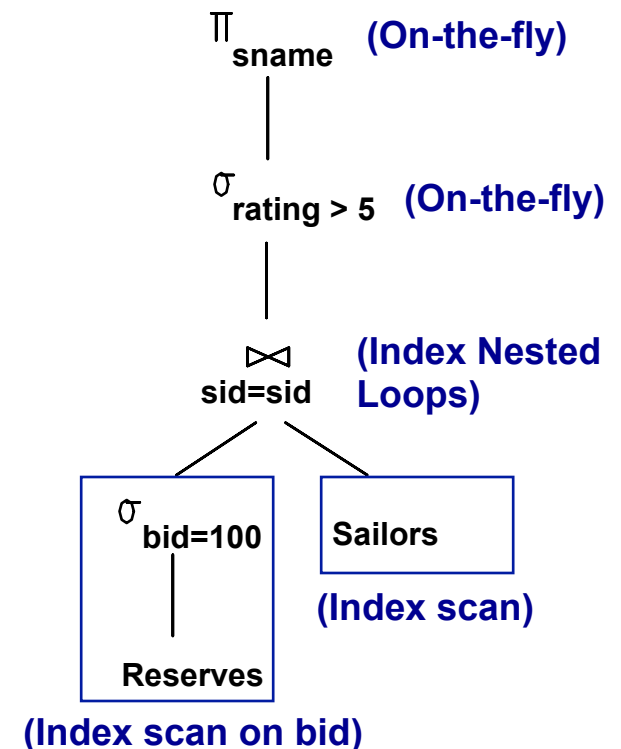
- Retrieve  $100,000/100 = 1000$  tuples
- Clustering: read  $1000/100 = 10$  pages.

- ❖ Indexed NLJ: *pipeline* the outer and *index lookup* on *sid* of Sailors.

- The outer: no need to materialize.
- The inner: *sid* is a *key*; *at most one* match tuple, unclustered index OK.

- ❖ Cost:

- Selection of Reserves tuples ( $\sim 10$  I/Os).
- For each tuple, get matching Sailor tuple ( $1000 \cdot (2 \sim 3)$ ).
- Total =  $2010 \sim 3010$  I/Os.





# Outline

---

- ❖ Query plans and equivalences

- ❖ Query optimization issues

- Plan space
- Cost estimation
- Plan search

- ❖ Handling nested queries

# Three Main Issues in Optimization

---

- ❖ Given a query block, three main optimization issues:
  - *Plan space*: what plans are considered?
  - *Plan cost*: what is the cost of a given plan?
  - *Search algorithm*: how do we search the plan space for the cheapest estimated plan?

# (1) Cost Estimation

- ❖ For each plan considered, must estimate its cost.
- ❖ Estimate *cost* of each operation in a plan tree:
  - Depends on input cardinalities.
  - Depends on the method (sequential scan, index scan, join...)
- ❖ Estimate *size of result* for each operation in tree:
  - Use statistics about input relations.
  - Estimate the *reduction factors* of predicates: *reduction factor (RF)* or *selectivity* of each *predicate* reflects the impact of the *term* in reducing result size.

```
SELECT attribute list  
FROM relation list  
WHERE term1 AND ... AND termk
```

# *Statistics in System Catalog*

---

- ❖ Statistics about each relation ( $R$ ) and index ( $I$ ):
  - Relation cardinality: # tuples (NTuples) in  $R$
  - Relation size: # pages (NPages) in  $R$
  - Index cardinality: # distinct values (NKeys) in  $I$
  - Index size: # pages (INPages) in  $I$
  - Index height: # nonleaf levels (IHeight) of  $I$
  - Index range: low/high key values (Low/High) in  $I$
  - Number of distinct values in an attribute (NKeys)
  - Histogram for an attribute

# Cost Estimates for Single-Relation Plans

- ❖ Index I on primary key matches selection:
  - Cost of lookup =  $\text{Height}(I)+1$  for a B+ tree,  $\approx 1.2$  for hash index
  - Cost of record retrieval = 1
- ❖ Clustered index I matching one or more selections:
  - Cost of lookup +  $\text{product of RF's of matching terms (RF-terms)} * (\text{INPages}'(I) + \text{NPages}(R))$
- ❖ Non-clustered index I matching one or more selections:
  - Cost of lookup +  $\text{RF-terms} * \text{INPages}'(I) + \min(\text{RF-terms} * \text{NTuples}(R), \text{NPages}(R))$
- ❖ Sequential scan of file:  $\text{NPages}(R)$
- ❖ May add extra costs for GROUP BY and duplicate elimination (if a query says DISTINCT)

# *Equiwidth Histograms*

**Equiwidth:** buckets of equal size.

<b>Frequency</b>	8/3	4/3	15/3	3/3	15/3										
<b>Counts</b>	8	4	15	3	15										
<b>Buckets</b>	<hr/>														
	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
	└──┘		└──┘		└──┘		└──┘		└──┘		└──┘		└──┘		


Still not accurate for  
value 14: 5/45

# Equidepth Histograms

**Equidepth**: equal counts of buckets.

Small errors for infrequent items: tolerable.

<b>Frequency</b>	9/4	10/4	10/2	7/4	9/1										
<b>Counts</b>	9	10	10	7	9										
<b>Buckets</b>	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14





Now accurate for value 14: 9/45

- ❖ Favors *frequent* values.
- ❖ Implementation:
  - Boundaries of 5 buckets {0, 4, 8, 10, 14, 14}
  - Count of tuples for each bucket
  - Number of distinct values for each bucket

## (2) Plan Space

- ❖ For each query block, the plans considered are:
  - All *access methods*, for each reln in the FROM clause.
  - All *left-deep join trees*: all the ways to join the relns one-at-a-time, with the inner reln in the FROM clause.
  - All permutations of N relns: **N factorial!**






# Plan Space

---

- ❖ For each block, the plans considered are:
  - All *access methods*, for each reln in FROM clause.
  - All *left-deep join trees*: all the ways to join the relns one-at-a-time, with the inner reln in the FROM clause.
    - All permutations of N relns: **N factorial!**
    - But avoid **Cartesian products!**  
Join R, S, T w.  $R.a = S.a$  and  $S.b = T.b$ , how many left-deep trees?
  - All *join methods*, for each join in the tree.
  - Appropriate *places for selections and projections*.

### (3) Plan Search


- ❖ As the number of joins increases, the number of alternative plans grows rapidly.
- ❖ System R: (1) use *only left-deep join trees*, (2) avoid *Cartesian products*.
  - Allow *pipelined* plans; intermediate results not written to temporary files.
  - Not all left-deep trees are fully pipelined!
    - Sort-Merge join: at least sorting phase
    - Two-phase hash join: partitioning phase



# Search Algorithm

## ❖ Left-deep join plans :

- Differ in the *order* of relations, *access method* for each relation, *join method* for each join.
- But may share common prefixes, so don't enumerate all. Instead use



## ❖ Dynamic Programming

“a method for solving problems that exhibit the properties of *overlapping subproblems* and *optimal substructures*”

- Find the best plans to access A, B, C, D individually
- Find the best plans for joining A-B, A-C, A-D, B-A, B-C, B-D, C-A, C-B, C-D, D-A, D-B, D-C; store the best for (A-B), (A-C), (A-D), (B-C), (B-D), (C-D)
- Repeat this for 3 relation sets...
- Repeat this for 4 relation sets...
- **This procedure is revised if given specific join predicates of A,B,C,D (i.e., left deep trees but avoiding Cartesian products).**

# An Example Star Schema

## ❖ Dynamic Programming

“a method for solving problems that exhibit the properties of *overlapping subproblems* and *optimal substructures*”

- Find the best plans to access A, B, C, D individually
- Find the best plans for joining A-B, A-C, A-D, B-A, C-A, D-A;; store the best for (A-B), (A-C), (A-D)
- Repeat this for 3 relation sets...
- Repeat this for 4 relation sets...

