

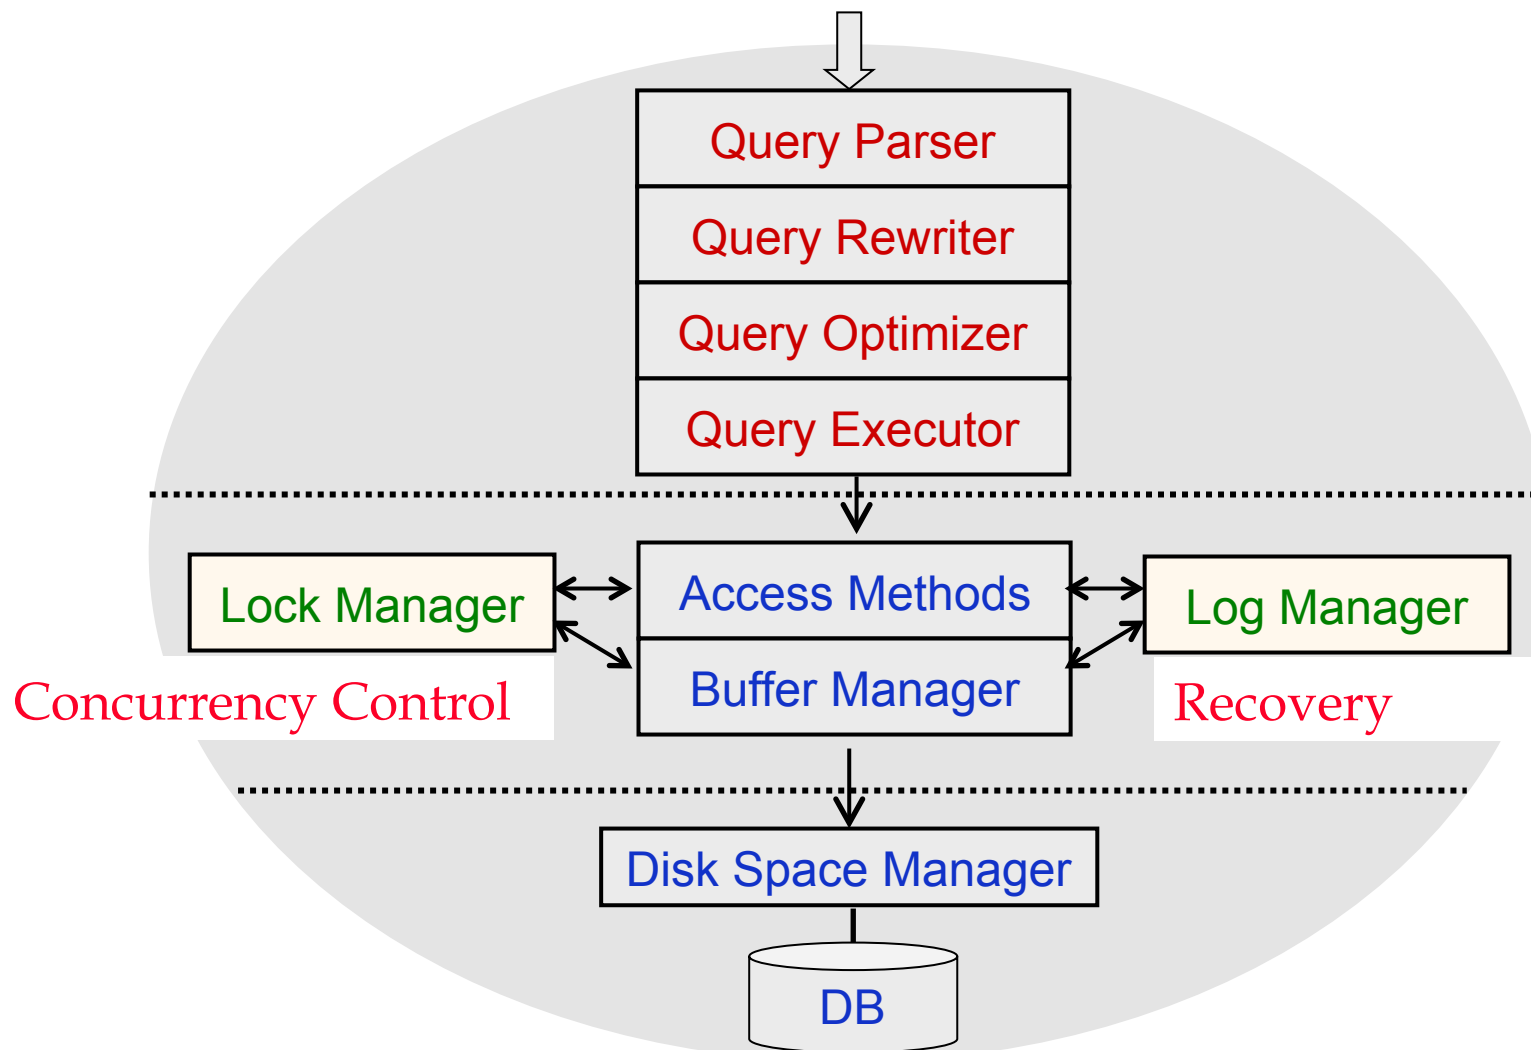


Transaction Management: Concurrency Control

Yanlei Diao



DBMS Architecture



Outline

- ❖ Transaction management overview
- ❖ Serializability & recoverability
- ❖ Lock-based concurrency control
- ❖ Optimistic concurrency control
- ❖ Efficient B+tree locking

Transactions

- ❖ User programs may do many things on the data retrieved.
 - E.g., operations on Bob's bank account.
 - E.g. transfer of money from account A to account B.
 - E.g., search for a ticket, think about it..., and buy it.
- ❖ But the DBMS is only concerned about what data is read from/written to the database.
- ❖ A *transaction* is DBMS's abstract view of a user program, simply, *a sequence of reads and writes*.

Concurrency

- ❖ Many users submit xacts, but each user thinks of his as executing by itself.
 - DMBS *interleaves* reads and writes of xacts for concurrency.
- ❖ Consistency: each xact starts and ends with a consistent state (i.e., satisfying all *integrity constraints*).
 - E.g., if an IC states that all accounts must have a positive balance, no transaction can violate this rule.
- ❖ Isolation: execution of one xact appears isolated from others.
 - Nobody else can see the data in its *intermediate state*, e.g., account A being debited but B not being credited.


Recovery

- ❖ A transaction might *commit* after completing all its actions, or it could be *aborted* after executing some actions.
- ❖ *Atomicity*: either all actions of a xact are performed or none of them is (*all-or-none*).
 - DBMS *logs* all actions so that it can *undo* the actions of aborted xacts.
- ❖ *Durability*: once a user program has been notified of success, its effect will persist despite system failure.
 - DBMS *logs* all actions so that it can *redo* the actions of committed xacts.

James Gray & Turing Award

❖ Jim Gray won Turing Award in 1998 for

“for seminal contributions to database and transaction processing research and technical leadership in system implementation”



Outline

- ❖ Transaction management overview
- ❖ Serializability & recoverability
- ❖ Lock-based concurrency control
- ❖ Optimistic concurrency control
- ❖ Efficient B+tree locking

Example

- ❖ Consider two transactions:

T1:	BEGIN	$A=A+100$,	$B=B-100$	END
T2:	BEGIN	$A=1.06*A$,	$B=1.06*B$	END

- 1st xact transfers \$100 from B' s account to A' s.
 - 2nd xact credits both accounts with a 6% interest payment.
 - No guarantee that T1 will execute before T2 or vice-versa, if both are submitted together.
-
- ❖ However, the net effect must be *equivalent to* these two transactions running *serially* in some order!

Example (Contd.)

- ❖ Consider a possible interleaving *schedule*:

T1:	$A=A+100,$	$B=B-100$
T2:	$A=1.06*A,$	$B=1.06*B$

- ❖ This is OK. But what about:

T1:	$A=A+100,$	$B=B-100$
T2:	$A=1.06*A, B=1.06*B$	

The DBMS' s view of the second schedule:

T1:	$R(A), W(A),$	$R(B), W(B)$
T2:	$R(A), W(A), R(B), W(B)$	

Scheduling Transactions

- ❖ Serial schedule: Schedule that does not interleave the actions of different transactions.
- ❖ Equivalent schedules: For any database state, the effect of executing the first schedule is identical to the effect of executing the second schedule.
- ❖ Serializable schedule: A schedule that is equivalent to some serial execution of the transactions.
 - If each transaction preserves consistency, every serializable schedule preserves consistency.

Serializability

- ❖ *Serializability theory* concerns the schedules of transactions that are not (explicitly) aborted.
- ❖ Given a set of such xacts, ideally want to allow *any serializable schedule*.
 - Recognizing any serializable schedule is highly complex, if possible.
- ❖ Instead, allow only a *subset* of serializable schedules that are easy to detect.

Conflict Serializability

- ❖ Two schedules are conflict equivalent if:
 - Involve the same actions of the same transactions.
 - Every pair of *conflicting actions* is ordered the same way.
- ❖ Schedule S is conflict serializable if S is conflict equivalent to some serial schedule.
- ❖ Given a set of xacts, conflict serializable schedules are a *subset* of serializable schedules.
 - There are serializable schedules that can't be detected using conflict serializability.

Dependency Graph


❖ Precedence graph:

- One node per Xact;
- Edge from Xact T_i to Xact T_j if an action of T_i precedes and conflicts with one of T_j 's actions (*RW, WR, WW operations on the same object*).

❖ **Theorem:** Schedule is conflict serializable *if and only if* its precedence graph is acyclic.

Example

T1:	R(A), W(A),	R(B), W(B)
T2:	R(A), W(A), R(B), W(B)	



- ❖ The schedule is not conflict serializable:
 - The cycle in the graph reveals the problem. The output of T1 depends on T2, and vice-versa.

Recoverability

- ❖ *Recoverability theory* concerns schedules that involve *aborted* transactions.

T1: R(A),W(A)	Abort
T2: R(A),W(A)	Commit

Unrecoverable!

- ❖ A schedule S is recoverable if each xact *commits* only after all xacts from which it read have committed.

Recoverability (Contd.)

T1: R(A),W(A)	Abort
T2: R(A),W(A)	Abort

Recoverable, but with cascading aborts.

- ❖ S avoids cascading rollback if each xact may *read* only those values written by committed xacts.

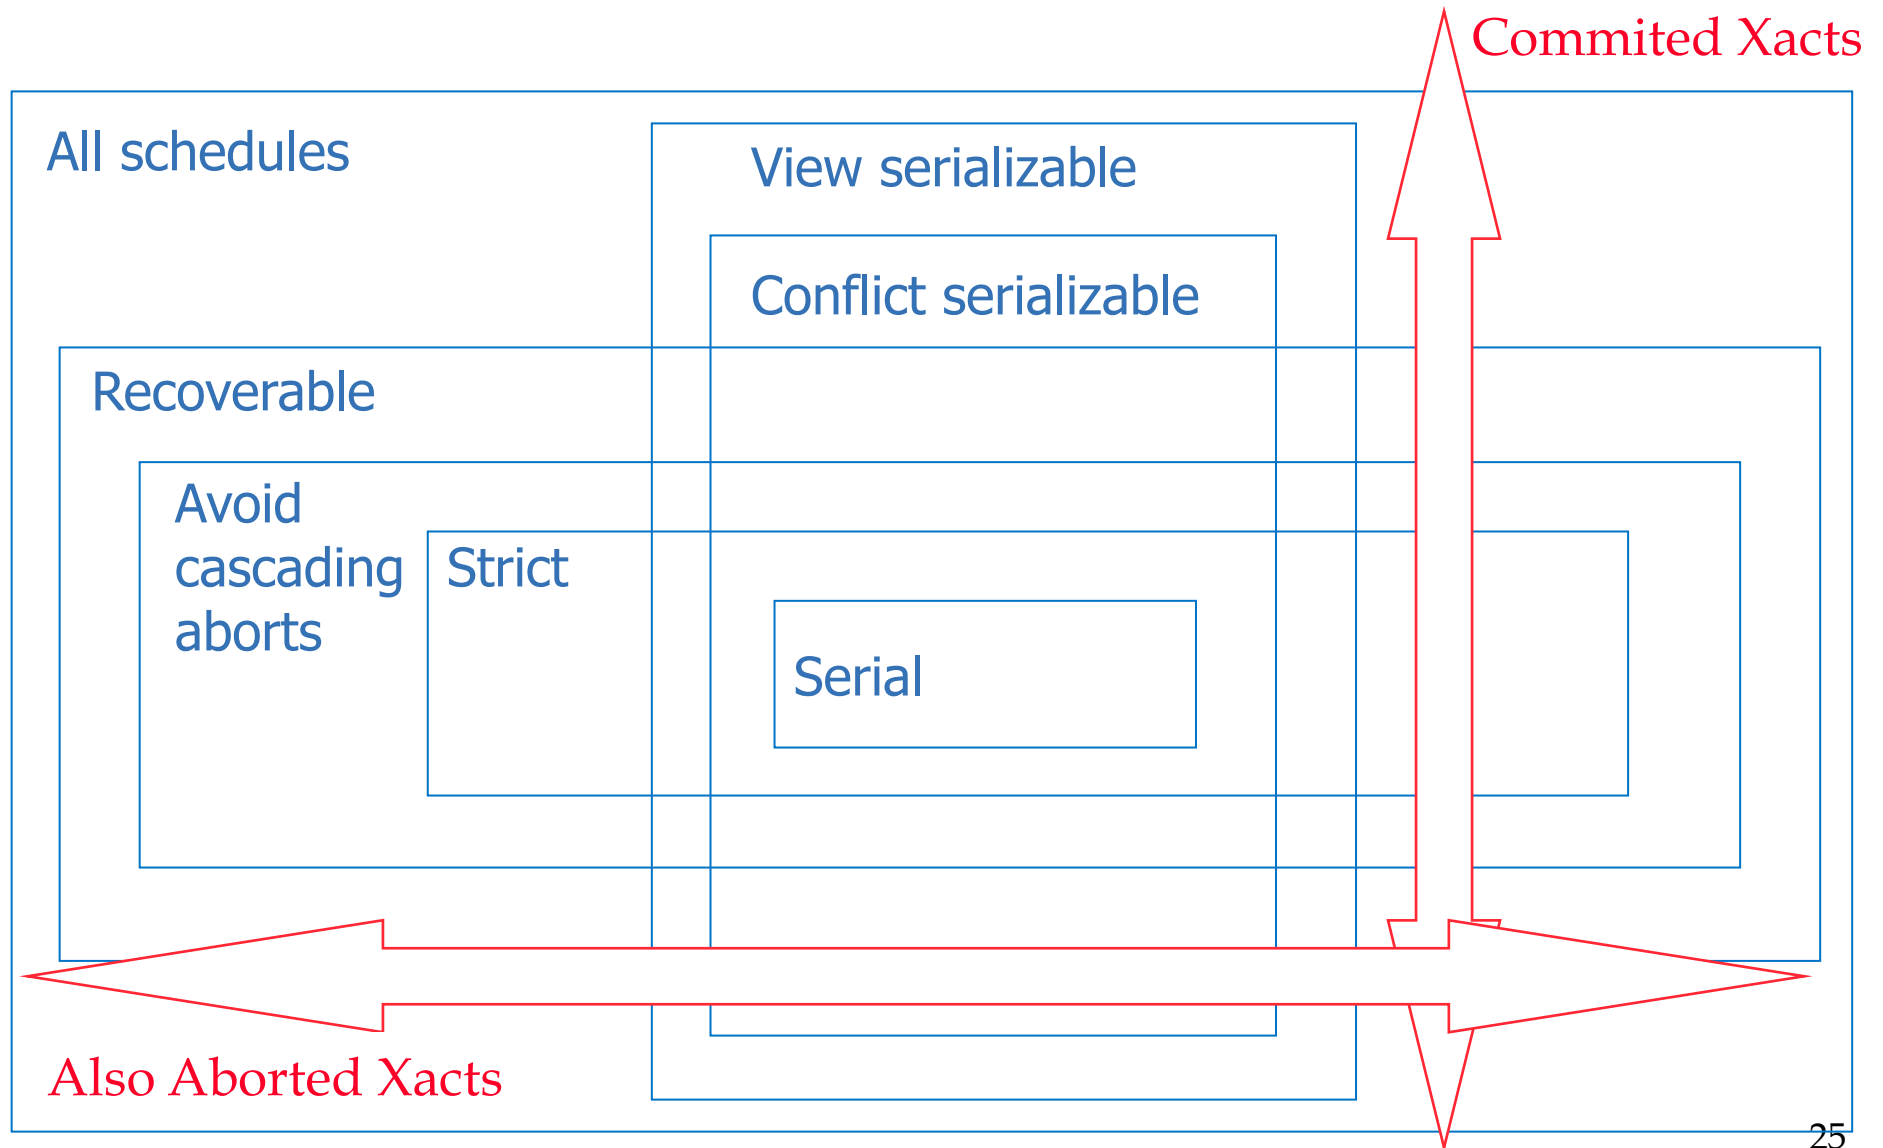
Recoverability (Contd.)

T1:	R(A), W(A)	Abort
T2:	R(A) W(A)	(Commit)

Recoverable, no cascading aborts,
but update of A by T2 will be lost!

- ❖ S is strict if each xact may *read and write* only objects *previously* written by committed xacts.
 - No cascading aborts.
 - Actions of aborted xacts can be simply undone by restoring the original values of modified objects.

Venn Diagram for Schedules



Outline


- ❖ Transaction management overview
- ❖ Serializability & recoverability
- ❖ Lock-based concurrency control
- ❖ Optimistic concurrency control
- ❖ Efficient B+tree locking

(1) Locking Protocol: Strict 2PL

❖ Strict Two-Phase Locking (Strict 2PL) Protocol:

1. Each Xact must obtain a **S (shared) lock** on object before reading, an **X (exclusive) lock** on object before writing.
2. If an Xact holds an X lock on an object, other Xact's cannot get a lock (S or X) on that object and become **blocked**.
3. All locks held by a xact are released when it completes.
 - Other blocked xact's can resume now.

Compatibility	Shared	Exclusive
Shared	Y	N
Exclusive	N	N




Strict 2PL (contd.)

- ❖ **Theorem:** Strict 2PL allows only schedules whose precedence graph is acyclic.
 - Strict 2PL only allows *conflict serializable* schedules!
- ❖ Strict 2PL is strict with respect to recoverability.
 - Strict 2PL is *recoverable without anomalies related to aborted transactions*.
 - Hence, it simplifies transaction aborts.

Locking Protocol: Nonstrict 2PL

❖ Nonstrict **Two-Phase** Locking Protocol

1. Each Xact must obtain a S (*shared*) lock on object before reading, an X (*exclusive*) lock on object before writing.
2. If an Xact holds an X lock on an object, other Xact's cannot get a lock (S or X) on that object and become *blocked*.
3. A xact *cannot request additional locks once it releases any locks.*
 - It releases locks earlier, so blocked xact's can resume earlier.




Nonstrict 2PL (contd.)

❖ **Theorem:** Nonstrict 2PL ensures acyclicity of precedence graph.


- Nonstrict 2PL only allows *conflict serializable* schedules.
- An equivalent serial schedule is given by the order of xacts entering their *shrinking phase*.

❖ Nonstrict 2PL is recoverable but *not strict*!

- Involves complex abort processing.
- But allows xacts to go through more quickly.



Deadlocks





- ❖ Deadlock: Cycle of transactions waiting for locks to be released by each other.
- ❖ Two ways of dealing with deadlocks:
 - Deadlock *detection*
 - Deadlock *prevention*

Deadlock Detection

- ❖ Create a **waits-for graph**:
 - Nodes are Xacts.
 - There is an edge from Xact T_i to Xact T_j if T_i is *waiting for* T_j to release a lock.
 - Note the difference from the precedence graph for conflict serializability.
- ❖ Periodically check for cycles, indicating deadlocks, in the waits-for graph.
 - Resolve a deadlock by aborting a transaction on the cycle and releasing all its locks.

Deadlock Detection (Contd.)

T1: S(A), R(A), S(B)
 T2: X(B), W(B) X(C)
 T3: S(C), R(C) X(B) **X(A)**
 T4:



Deadlock Prevention

- ❖ Assign priorities based on timestamps.
 - The older the timestamp, the higher the xact's priority.
- ❖ Wait-Die: T_i wants a lock that T_j holds. If T_i has higher priority, T_i waits for T_j ; otherwise T_i aborts.
 - *Lower priority xacts can never wait.*
- ❖ Wound-wait: T_i wants a lock that T_j holds. If T_i has higher priority, T_j aborts; otherwise T_i waits.
 - *Higher priority xacts never wait.*
- ❖ If a transaction re-starts, make sure it has its original timestamp so its priority increases.